

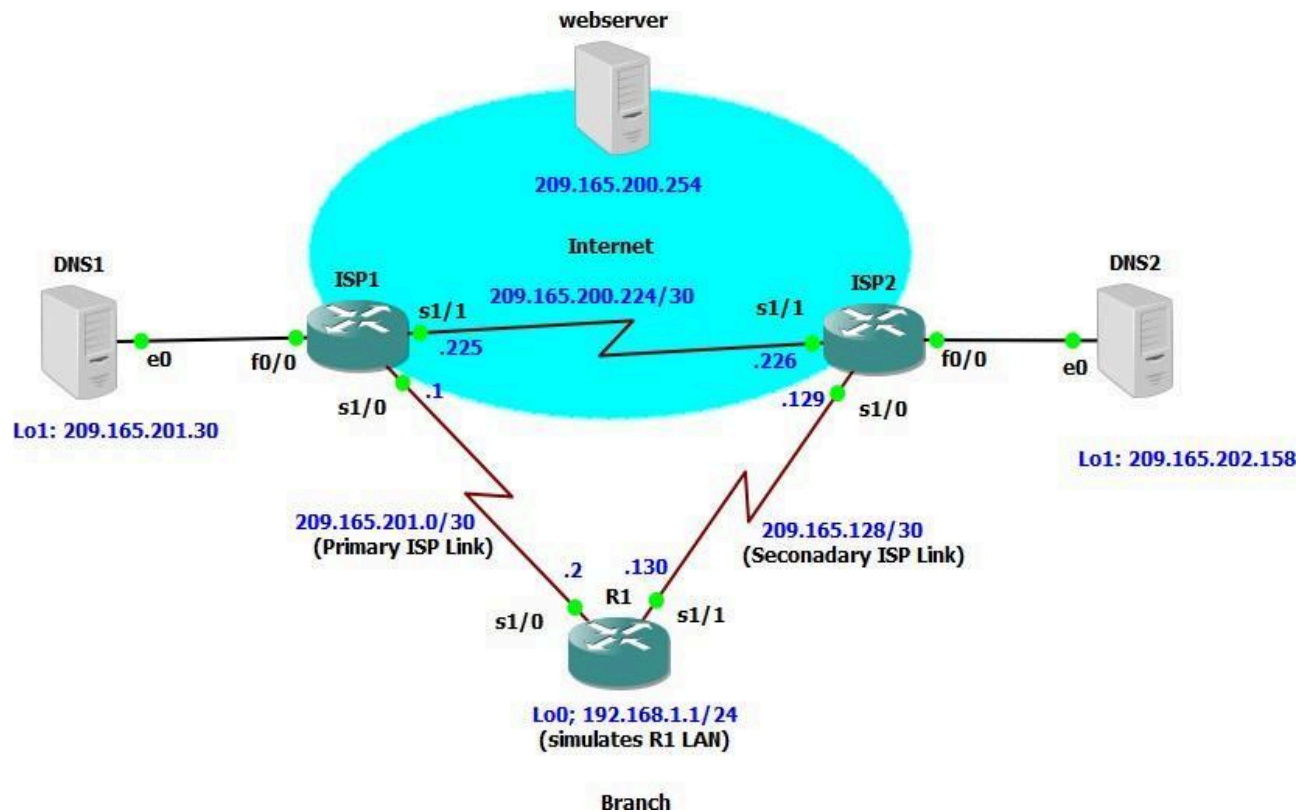
INDEX

Prac. No.	Practical	Date
1	Configure IP SLA Tracking and Path Control Topology	
2	Using the AS_PATH Attribute	
3	Configuring IBGP and EBGP Sessions, Local Preference, and MED	
4	Secure the Management Plane	
5	Configure and Verify Path Control Using PBR s	
6	Simulating MPLS environment	

Practical No - 1

Aim: Configure IP SLA Tracking and Path Control Topology

Topology :



Objectives

- Configure and verify the IP SLA feature.
- Test the IP SLA tracking feature.
- Verify the configuration and operation using show and debug commands.

Step 1: Prepare the routers and configure the router hostname and interface addresses.

Router R1

```
interface Loopback 0
```

```
ip address 192.168.1.1 255.255.255.0
```

```
interface Serial0/0/0
```

```
ip address 209.165.201.2 255.255.255.252
```

```
no shutdown
```

```
interface
```

```
Serial0/0/1
```

```
ip address 209.165.202.130 255.255.255.252
```

```
no shutdown
```

```
R1(config)#interface Loopback 0
```

```
R1(config-if)#ip address 192.168.1.1 255.255.255.0
```

```
R1(config-if)#
```

```
R1(config-if)#int s1/0
```

```
R1(config-if)#ip address 209.165.201.2 255.255.255.252
```

```
R1(config-if)#no shutdown
```

```
R1(config-if)#int s1/1
```

```
R1(config-if)#ip address 209.165.202.130 255.255.255.252
```

```
R1(config-if)#no shutdown
```

Router ISP1 (R2)

```
interface Loopback0
```

```
ip address 209.165.200.254 255.255.255.255
```

```
interface Loopback1
```

```
ip address 209.165.201.30 255.255.255.255
```

```
interface Serial0/0/0
```

```
ip address 209.165.201.1 255.255.255.252
```

```
no shutdown
```

```
interface
```

```
Serial0/0/1
```

```
ip address 209.165.200.225 255.255.255.252
```

```
no shutdown
```

```

ISP1(config)#interface Loopback0
ISP1(config-if)#
*May 18 15:24:24.315: %LINEPROTO-5-UPDOWN: Line protocol on
ISP1(config-if)#ip address 209.165.200.254 255.255.255.255
ISP1(config-if)#interface Loopback1
ISP1(config-if)#
*May 18 15:24:36.915: %LINEPROTO-5-UPDOWN: Line protocol on
ISP1(config-if)#ip address 209.165.201.30 255.255.255.255
ISP1(config-if)#int s1/0
ISP1(config-if)#ip address 209.165.201.1 255.255.255.252
ISP1(config-if)#no shutdown
ISP1(config-if)#
*May 18 15:25:03.695: %LINK-3-UPDOWN: Interface Serial1/0,
ISP1(config-if)#
ISP1(config-if)#i
*May 18 15:25:04.699: %LINEPROTO-5-UPDOWN: Line protocol on
ISP1(config-if)#int s1/1
ISP1(config-if)#ip address 209.165.200.225 255.255.255.252
ISP1(config-if)#no shutdown

```

Router ISP2 (R3)

```

interface Loopback0
ip address 209.165.200.254 255.255.255.255
interface Loopback1
ip address 209.165.202.158 255.255.255.255
interface Serial0/0/0
description ISP2 -->
R1
ip address 209.165.202.129 255.255.255.252
no shutdown
interface
Serial0/0/1
ip address 209.165.200.226 255.255.255.252
no shutdown

```

```

ISP2(config)#interface Loopback0
ISP2(config-if)#
*May 18 15:25:22.219: %LINEPROTO-5-UPDOWN: Line protocol on
ISP2(config-if)#ip address 209.165.200.254 255.255.255.255
ISP2(config-if)#interface Loopback1
ISP2(config-if)#
*May 18 15:25:34.595: %LINEPROTO-5-UPDOWN: Line protocol on
ISP2(config-if)#ip address 209.165.202.158 255.255.255.255
ISP2(config-if)#int s1/0
ISP2(config-if)#ip address 209.165.202.129 255.255.255.252
ISP2(config-if)#no shutdown
ISP2(config-if)#
ISP2(config-if)#
*May 18 15:26:01.299: %LINK-3-UPDOWN: Interface Serial1/0,
ISP2(config-if)#i
*May 18 15:26:02.303: %LINEPROTO-5-UPDOWN: Line protocol on
ISP2(config-if)#int s1/1
ISP2(config-if)#ip address 209.165.200.226 255.255.255.252
ISP2(config-if)#no shutdown

```

b. Verify the configuration by using the show interfaces description command. The output from router R1 is shown here as an example.

R1# show interfaces description

```

R1#show interfaces description
Interface          Status      Protocol Description
Fa0/0              admin down  down
Se1/0              up          up
Se1/1              up          up
Se1/2              admin down  down
Se1/3              admin down  down
Lo0                up          up

```

c. The current routing policy in the topology is as follows:

- Router R1 establishes connectivity to the Internet through ISP1 using a default static route.
- ISP1 and ISP2 have dynamic routing enabled between them, advertising their respective public address pools.
- ISP1 and ISP2 both have static routes back to the ISP LAN.

Router R1

ip route 0.0.0.0 0.0.0.0 209.165.201.1

```

R1(config)#
R1(config)# ip route 0.0.0.0 0.0.0.0 209.165.201.1

```

Router ISP1 (R2)

router eigrp 1

```
network 209.165.200.224 0.0.0.3
```

```
network 209.165.201.0 0.0.0.31
```

```
no auto-summary
```

```
ip route 192.168.1.0 255.255.255.0 209.165.201.2
```

```
ISP1(config)#router eigrp 1
ISP1(config-router)#network 209.165.200.224 0.0.0.3
ISP1(config-router)#network 209.165.201.0 0.0.0.31
ISP1(config-router)#no auto-summary
ISP1(config-router)#ip route 192.168.1.0 255.255.255.0 209.165.201.2
```

Router ISP2 (R3)

```
router eigrp 1
```

```
network 209.165.200.224 0.0.0.3
```

```
network 209.165.202.128 0.0.0.31
```

```
no auto-summary
```

```
ip route 192.168.1.0 255.255.255.0 209.165.202.130
```

```
ISP2(config)#router eigrp 1
ISP2(config-router)#network 209.165.200.224 0.0.0.3
ISP2(config-router)#
*May 18 15:30:14.515: %DUAL-5-NBRCHANGE: IP-EIGRP(0) 1: Neighbor 209.165.201.2 (FastEthernet0/0) is up: 0:0:0
ISP2(config-router)#network 209.165.202.128 0.0.0.31
ISP2(config-router)#no auto-summary
ISP2(config-router)#
*May 18 15:30:28.971: %DUAL-5-NBRCHANGE: IP-EIGRP(0) 1: Neighbor 209.165.201.2 (FastEthernet0/0) is up: 0:0:0
ISP2(config-router)#ip route 192.168.1.0 255.255.255.0 209.165.202.130
```

Step 2: Verify server reachability.

- a. Before implementing the Cisco IOS SLA feature, you must verify reachability to the Internet servers. From router R1, ping the web server, ISP1 DNS server, and ISP2 DNS server to verify connectivity. You can copy the following Tcl script and paste it into R1.

```
R1(tcl)# foreach address {
```

```
+>(tcl)# 209.165.200.254
```

```
+>(tcl)# 209.165.201.30
```

```
+>(tcl)# 209.165.202.158
```

```
+>(tcl)# } { ping $address source 192.168.1.1 }
```

```
R1#tclsh
R1(tcl)#foreach address {
+>(tcl)#209.165.200.254
+>(tcl)#209.165.201.30
+>(tcl)#209.165.202.158
+>(tcl)#} {
+>(tcl)#ping $address source 192.168.1.1
+>(tcl)#}
```

```
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 209.165.200.254, timeout is 2 seconds:
Packet sent with a source address of 192.168.1.1
!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 28/78/96 ms
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 209.165.201.30, timeout is 2 seconds:
Packet sent with a source address of 192.168.1.1
!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 20/31/48 ms
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 209.165.202.158, timeout is 2 seconds:
Packet sent with a source address of 192.168.1.1
!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 16/37/60 ms
```

b. Trace the path taken to the web server, ISP1 DNS server, and ISP2 DNS server. You can copy the following Tcl script and paste it into R1.

```
R1(tcl)# foreach address {
+>(tcl)# 209.165.200.254
+>(tcl)# 209.165.201.30
+>(tcl)# 209.165.202.158
+>(tcl)# } { trace $address source 192.168.1.1 }
```

```
R1(tcl)#foreach address {
+>(tcl)#209.165.200.254
+>(tcl)#209.165.201.30
+>(tcl)#209.165.202.158
+>(tcl)#} {
+>(tcl)#trace $address source 192.168.1.1
+>(tcl)#}
```

```

Type escape sequence to abort.
Tracing the route to 209.165.200.254

  1 209.165.201.1 24 msec 24 msec 16 msec
Type escape sequence to abort.
Tracing the route to 209.165.201.30

  1 209.165.201.1 16 msec 24 msec 24 msec
Type escape sequence to abort.
Tracing the route to 209.165.202.158

  1 209.165.201.1 24 msec 24 msec 32 msec
  2 209.165.200.226 28 msec 44 msec 72 msec

```

Step 3: Configure IP SLA probes.

- a. Create an ICMP echo probe on R1 to the primary DNS server on ISP1 using the ip sla command. the previous ip sla monitor command. In addition, the icmp-echo command has replaced the type echo protocol ipIcmpEcho command.

```
R1(config)# ip sla 11
```

```
R1(config-ip-sla)# icmp-echo
```

```
209.165.201.30 R1(config-ip-sla-echo)#
```

```
frequency 10 R1(config-ip-sla-echo)# exit
```

```
R1(config)# ip sla schedule 11 life forever start-time now
```

```

R1(config)#ip sla 11
R1(config-ip-sla)#icmp-echo 209.165.201.30
R1(config-ip-sla-echo)#frequency 10
R1(config-ip-sla-echo)#exit
R1(config)#ip sla schedule 11 life forever start-time now
R1(config)#exit

```

- b. Verify the IP SLAs configuration of operation 11 using the show ip sla configuration 11 command.

```
R1# show ip sla configuration 11
```



```

R1(tcl)#show ip sla configuration 11
IP SLAs, Infrastructure Engine-II.
Entry number: 11
Owner:
Tag:
Type of operation to perform: icmp-echo
Target address/Source address: 209.165.201.30/0.0.0.0
Type Of Service parameter: 0x0
Request size (ARR data portion): 28
Operation timeout (milliseconds): 5000
Verify data: No
Vrf Name:
Schedule:
  Operation frequency (seconds): 10 (not considered if
  Next Scheduled Start Time: Start Time already passed
  Group Scheduled : FALSE
  Randomly Scheduled : FALSE
  Life (seconds): Forever
  Entry Ageout (seconds): never
  Recurring (Starting Everyday): FALSE
  Status of entry (SNMP RowStatus): Active

```

- c. Issue the show ip sla statistics command to display the number of successes, failures, and results of the latest operations.

R1# show ip sla statistics

```

R1#show ip sla statistics 22
IPSLAs Latest Operation Statistics

IPSLA operation id: 22
Type of operation: icmp-echo
  Latest RTT: 64 milliseconds
Latest operation start time: *15:40:40.823 UTC Tue May 18 2021
Latest operation return code: OK
Number of successes: 6
Number of failures: 0
Operation time to live: Forever

```

- d. Although not actually required because IP SLA session 11 alone could provide the desired fault tolerance, create a second probe, 22, to test connectivity to the second DNS server located on router ISP2. You can copy and paste the following commands on R1.

```

R1(config)#
R1(config)#ip sla 22
R1(config-ip-sla)#icmp-echo 209.165.202.158
R1(config-ip-sla-echo)#frequency 10
R1(config-ip-sla-echo)#exit
R1(config)#ip sla schedule 22 life forever start-time now

```

- e. Verify the new probe using the show ip sla configuration and show ip sla statistics commands.

R1# show ip sla configuration 22

```

R1#show ip sla configuration 22
IP SLAs, Infrastructure Engine-II.
Entry number: 22
Owner:
Tag:
Type of operation to perform: icmp-echo
Target address/Source address: 209.165.202.158/0.0.0.0
Type Of Service parameter: 0x0
Request size (ARR data portion): 28
Operation timeout (milliseconds): 5000
Verify data: No

```

R1# show ip sla statistics 22

```

R1#show ip sla statistics 22
IPSLAs Latest Operation Statistics

IPSLA operation id: 22
Type of operation: icmp-echo
    Latest RTT: 64 milliseconds
Latest operation start time: *15:40:40.823 UTC Tue May 18 2021
Latest operation return code: OK
Number of successes: 6
Number of failures: 0
Operation time to live: Forever

```

Step 4: Configure tracking options.

- a. Remove the current default route on R1, and replace it with a floating static route having an administrative distance of 5.

```
R1(config)# no ip route 0.0.0.0 0.0.0.0 209.165.201.1
```

```
R1(config)# ip route 0.0.0.0 0.0.0.0 209.165.201.1 5
```

```
R1(config)# exit
```

```

R1(config)#
R1(config)#no ip route 0.0.0.0 0.0.0.0 209.165.201.1
R1(config)#ip route 0.0.0.0 0.0.0.0 209.165.201.1 5
R1(config)#exit

```

- b. Verify the routing

table. R1# show ip route

```

R1#show ip route
Codes: C - connected, S - static, R - RIP, M - mobile, B - BGP
       D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter area
       N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type 2
       E1 - OSPF external type 1, E2 - OSPF external type 2
       i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1, L2 - IS-IS level-2
       ia - IS-IS inter area, * - candidate default, U - per-user static route
       o - ODR, P - periodic downloaded static route

Gateway of last resort is 209.165.201.1 to network 0.0.0.0

```

- c. Use the track 1 ip sla 11 reachability command to enter the config-track subconfiguration mode.

```
R1(config)# track 1 ip sla 11 reachability
```

```
R1(config-track)#
```

```
R1(config)#
R1(config)#track 1 ip sla 11 reachability
R1(config-track)#delay down 10 up 1
R1(config-track)#exit
```

- d. Configure the floating static route that will be implemented when tracking object 1 is active. To view routing table changes as they happen, first enable the debug ip routing command. Next, use the ip route 0.0.0.0 0.0.0.0 209.165.201.1 2 track 1 command to create a floating static default route via 209.165.201.1 (ISP1). Notice that this command references the tracking object number 1, which in turn references IP SLA operation number 11.

```
R1# debug ip routing
```

```
R1#debug ip routing
IP routing debugging is on
```

```
R1(config)# ip route 0.0.0.0 0.0.0.0 209.165.201.1 2 track 1
```

```
R1(config)#ip route 0.0.0.0 0.0.0.0 209.165.201.1 2 track 1
R1(config)#
*May 18 15:43:00.035: RT: closer admin distance for 0.0.0.0, flushing 1 routes
*May 18 15:43:00.035: RT: NET-RED 0.0.0.0/0
*May 18 15:43:00.035: RT: add 0.0.0.0/0 via 209.165.201.1, static metric [2/0]
*May 18 15:43:00.039: RT: NET-RED 0.0.0.0/0
*May 18 15:43:00.039: RT: default path is now 0.0.0.0 via 209.165.201.1
*May 18 15:43:00.039: RT: new default network 0.0.0.0
*May 18 15:43:00.043: RT: NET-RED 0.0.0.0/0
```

- e. Repeat the steps for operation 22, track number 2, and assign the static route an admin distance higher than track 1 and lower than 5. On R1, copy the following configuration, which sets an admin distance of 3. track 2 ip sla 22 reachability delay down 10 up 1 exit

```
ip route 0.0.0.0 0.0.0.0 209.165.202.129 3 track 2
```

```
R1(config)#track 1 ip sla 22 reachability
R1(config-track)#delay down 10 up 1
R1(config-track)#exi
*May 18 15:43:56.339: RT: NET-RED 0.0.0.0/0
R1(config-track)#exit
R1(config)#ip route 0.0.0.0 0.0.0.0 209.165.201.1 2 track 1
R1(config)#ip route 0.0.0.0 0.0.0.0 209.165.202.129 3 track 2
R1(config)#exit
```

- f. Verify the routing table again.

R1# show ip route

```
R1#show ip route
Codes: C - connected, S - static, R - RIP, M - mobile, B - B
       D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF i
       N1 - OSPF NSSA external type 1, N2 - OSPF NSSA extern
       E1 - OSPF external type 1, E2 - OSPF external type 2
       i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1, L2
       ia - IS-IS inter area, * - candidate default, U - per
       o - ODR, P - periodic downloaded static route

Gateway of last resort is 209.165.201.1 to network 0.0.0.0

    209.165.201.0/30 is subnetted, 1 subnets
C       209.165.201.0 is directly connected, Serial1/0
    209.165.202.0/30 is subnetted, 1 subnets
C       209.165.202.128 is directly connected, Serial1/1
C       192.168.1.0/24 is directly connected, Loopback0
S*    0.0.0.0/0 [2/0] via 209.165.201.1
```

Step 5: Verify IP SLA operation.

The following summarizes the process:

- Disable the DNS loopback interface on ISP1 (R2).
- Observe the output of the debug command on R1.
- Verify the static route entries in the routing table and the IP SLA statistics of R1.
- Re-enable the loopback interface on ISP1 (R2) and again observe the operation of the IP SLA tracking feature.

ISP1(config)# interface loopback

1 ISP1(config-if)# shutdown

```
ISP1(config)#
ISP1(config)#interface loopback 1
ISP1(config-if)#shutdown
```

b. Verify the routing

table. R1# show ip route

```

R1#show ip route
Codes: C - connected, S - static, R - RIP, M - mobile, B - BGP
        D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter area
        N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type 2
        E1 - OSPF external type 1, E2 - OSPF external type 2, O - OSPF other
        i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1, L2 - IS-IS level-2
        ia - IS-IS inter area, * - candidate default, U - per-user static route
        o - ODR, P - periodic downloaded static route

Gateway of last resort is 209.165.201.1 to network 0.0.0.0

    209.165.201.0/30 is subnetted, 1 subnets
C       209.165.201.0 is directly connected, Serial1/0
    209.165.202.0/30 is subnetted, 1 subnets
C       209.165.202.128 is directly connected, Serial1/1
C       192.168.1.0/24 is directly connected, Loopback0
S*    0.0.0.0/0 [2/0] via 209.165.201.1

```

c. Verify the IP SLA statistics.

R1# show ip sla statistics

```

R1#show ip sla statistics
IPSLAs Latest Operation Statistics

IPSLA operation id: 11
Type of operation: icmp-echo
    Latest RTT: NoConnection/Busy/Timeout
Latest operation start time: *15:47:41.887 UTC Tue May 18 2021
Latest operation return code: No connection
Number of successes: 51
Number of failures: 13
Operation time to live: Forever

```

d. Initiate a trace to the web server from the internal LAN IP address.

R1# trace 209.165.200.254 source 192.168.1.1

```

R1#trace 209.165.200.254 source 192.168.1.1

Type escape sequence to abort.
Tracing the route to 209.165.200.254

  1 209.165.201.1 4 msec 32 msec 32 msec

```

f. Again examine the IP SLA statistics.

R1# show ip sla statistics

```
R1#show ip sla statistics
IPSLAs Latest Operation Statistics

IPSLA operation id: 11
Type of operation: icmp-echo
    Latest RTT: 57 milliseconds
Latest operation start time: *15:50:01.887 UTC Tue May 18 2021
Latest operation return code: OK
Number of successes: 61
Number of failures: 17
Operation time to live: Forever
```

g. Verify the routing table.

R1# show ip route

```
R1#show ip route
Codes: C - connected, S - static, R - RIP, M - mobile, B - BGP
       D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter area
       N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type 2
       E1 - OSPF external type 1, E2 - OSPF external type 2
       i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1, L2 - IS-IS level-2
       ia - IS-IS inter area, * - candidate default, U - per-user static route
       o - ODR, P - periodic downloaded static route

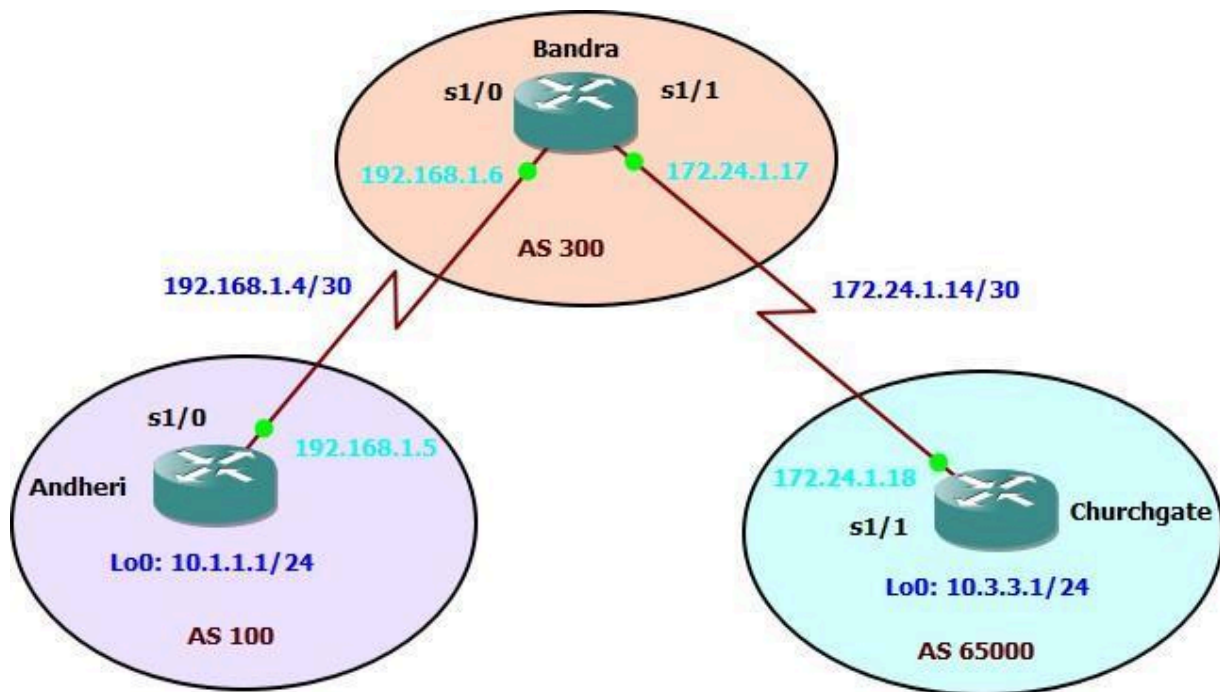
Gateway of last resort is 209.165.201.1 to network 0.0.0.0

    209.165.201.0/30 is subnetted, 1 subnets
C       209.165.201.0 is directly connected, Serial1/0
    209.165.202.0/30 is subnetted, 1 subnets
C       209.165.202.128 is directly connected, Serial1/1
C       192.168.1.0/24 is directly connected, Loopback0
S*     0.0.0.0/0 [2/0] via 209.165.201.1
```

Practical No - 2

Aim: Using the AS_PATH Attribute

Topolgy :



Objective :

- Use BGP commands to prevent private AS numbers from being advertised to the outside world.
- Use the AS_PATH attribute to filter BGP routes based on their source AS number

Step 1 : Prepare the routers for the lab.

Cable the network as shown in the topology diagram. Erase the startup configuration and reload each router to clear previous configurations.

Step 2 : Configure the hostname and interface addresses.

a. You can copy and paste the following configurations into your routers to begin.

Router R1 (hostname Andheri)

Andheri(config)# interface Loopback0

Andheri(config-if)# ip address 10.1.1.1 255.255.255.0

Andheri(config-if)# exit

Andheri(config)# interface Serial0/0/0

Andheri(config-if)# ip address 192.168.1.5 255.255.255.252

Andheri(config-if)# no shutdown

Andheri(config-if)# end

Andheri#

```
R1#
R1#conf t
Enter configuration commands, one per line. End with CNTL/Z.
R1(config)#hostname Andheri
Andheri(config)#int loopback 0
Andheri(config-if)#
*May  7 09:30:42.867: %LINEPROTO-5-UPDOWN: Line protocol on Interface Loopback0
Andheri(config-if)#ip address 10.1.1.1 255.255.255.0
Andheri(config-if)#exit
Andheri(config)#int s1/0
Andheri(config-if)#ip address 192.168.1.5 255.255.255.252
Andheri(config-if)#no shutdown
Andheri(config-if)#
*May  7 09:31:41.315: %LINK-3-UPDOWN: Interface Serial1/0, changed state to up
```

Router R2 (hostname Bandra)

Bandra(config)# interface Loopback0

Bandra(config-if)# ip address 10.2.2.1 255.255.255.0

Bandra(config-if)# interface Serial0/0/0

Bandra(config-if)# ip address 192.168.1.6 255.255.255.252

Bandra(config-if)# no shutdown

Bandra(config-if)# exit

Bandra(config)# interface Serial0/0/1

Bandra(config-if)# ip address 172.24.1.17 255.255.255.252

Bandra(config-if)# no shutdown

Bandra(config-if)#

end Bandra#


```

R2#
R2#conf t
Enter configuration commands, one per line. End with CNTL/Z
R2(config)#hostname Bandra
Bandra(config)#int loopback 0
Bandra(config-if)#ip addr
*May 7 09:31:30.407: %LINEPROTO-5-UPDOWN: Line protocol on Interface Loopback0,
Bandra(config-if)#ip address 10.2.2.1 255.255.255.0
Bandra(config-if)#exit
Bandra(config)#int s1/0
Bandra(config-if)#ip address 192.168.1.6 255.255.255.252
Bandra(config-if)#no shutdown
Bandra(config-if)#exit

```

Router R3 (hostname ChurchGate)

```
Churchgate(config)# interface Loopback0
```

```
Churchgate(config-if)# ip address 10.3.3.1
```

```
255.255.255.0 Churchgate(config-if)# exit
```

```
Churchgate(config)# interface Serial0/0/1
```

```
Churchgate(config-if)# ip address 172.24.1.18
```

```
255.255.255.252 Churchgate(config-if)# no shutdown
```

```
Churchgate(config-if)# end
```

```
Churchgate#
```

```

Bandra(config)#int s1/1
Bandra(config-if)#ip address 172.24.1.17 255.255.255.252
Bandra(config-if)#no shutdown
Bandra(config-if)#
*May 7 09:33:39.591: %LINK-3-UPDOWN: Interface Serial1/1,
Bandra(config-if)#

```

```

R3#
R3#conf t
Enter configuration commands, one per line. End with CNTL/Z.
R3(config)#hostname Churchgate
Churchgate(config)#int loopback 0
Churchgate(config-if)#ip[
*May 7 09:33:31.243: %LINEPROTO-5-UPDOWN: Line protocol on Interface Loopback0,
Churchgate(config-if)#ip address 10.3.3.1 255.255.255.0
Churchgate(config-if)#exit
Churchgate(config)#int s1/1
Churchgate(config-if)#ip address 172.24.1.18 255.255.255.252
Churchgate(config-if)#no shutdown
Churchgate(config-if)#
*May 7 09:34:39.795: %LINK-3-UPDOWN: Interface Serial1/1, ch

```

b. Use ping to test the connectivity between the directly connected routers.

```

Bandra#ping 192.168.1.5

Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 192.168.1.5, timeout is 2 seconds:
!!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 20/32/40 ms
Bandra#
Bandra#
Bandra#ping 172.24.1.18

Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 172.24.1.18, timeout is 2 seconds:
!!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 28/36/48 ms
Bandra#

```

Step 3 : Configure BGP.

- a. Configure BGP for normal operation. Enter the appropriate BGP commands on each router so that they identify their BGP neighbors and advertise their loopback networks.

```
Andheri(config)# router bgp 100
```

```
Andheri(config-router)# neighbor 192.168.1.6 remote-as 300
```

```
Andheri(config-router)# network 10.1.1.0 mask
```

```
255.255.255.0
```

```

Andheri#conf t
Enter configuration commands, one per line. End with CNTL
Andheri(config)#router bgp 100
Andheri(config-router)#neighbor 192.168.1.6 remote-as 300
Andheri(config-router)#network 10.1.1.0 mask 255.255.255.0

```

```
Bandra(config)# router bgp 300
```

```
Bandra(config-router)# neighbor 192.168.1.5 remote-as 100
```

```
Bandra(config-router)# neighbor 172.24.1.18 remote-as
```

```
65000 Bandra(config-router)# network 10.2.2.0 mask
```

```
255.255.255.0
```

```

Bandra#conf t
Enter configuration commands, one per line. End with CNTL
Bandra(config)#router bgp 300
Bandra(config-router)#neighbor 192.168.1.5 remote-as 100
Bandra(config-router)#
*May 7 10:04:59.051: %BGP-5-ADJCHANGE: neighbor 192.168.1.
Bandra(config-router)#neighbor 172.24.1.18 remote-as 65000
Bandra(config-router)#network 10.2.2.0 mask 255.255.255.0

```

```
Churchgate(config)# router bgp 65000
```

```
Churchgate(config-router)# neighbor 172.24.1.17 remote-as 300
```

```
Churchgate(config-router)# network 10.3.3.0 mask
```

```
255.255.255.0
```

```
Churchgate#conf t
Enter configuration commands, one per line. End with CNTL/Z.
Churchgate(config)#router bgp 65000
Churchgate(config-router)#neighbor 172.24.1.17 remote-as 300
Churchgate(config-router)#
*May  7 10:04:44.195: %BGP-5-ADJCHANGE: neighbor 172.24.1.17 U
Churchgate(config-router)#network 10.3.3.0 mask 255.255.255.0
```

- b. Verify that these routers have established the appropriate neighbor relationships by issuing the show ip bgp neighbors command on each router.

Bandra# show ip bgp neighbors

```
Bandra#show ip bgp neighbors
BGP neighbor is 172.24.1.18, remote AS 65000, external link
  BGP version 4, remote router ID 10.3.3.1
  BGP state = Established, up for 00:01:30
  Last read 00:00:13, last write 00:00:44, hold time is 180, keepalive
  seconds
  Neighbor capabilities:
    Route refresh: advertised and received(new)
    New ASN Capability: advertised and received
    Address family IPv4 Unicast: advertised and received
  Message statistics:
    InQ depth is 0
    OutQ depth is 0

              Sent          Rcvd
Opens:             1           1
Notifications:     0           0
Updates:           3           1
Keepalives:        2           3
Route Refresh:     0           0
Total:             6           5
Default minimum time between advertisement runs is 30 seconds
```

Step 4 : Remove the private AS.

- a. DBandralay the Andheri routing table using the show ip route command. Andheri should have a route to both 10.2.2.0 and 10.3.3.0. Troubleshoot if necessary.

Andheri#show ip route

```

Andheri#show ip route
Codes: C - connected, S - static, R - RIP, M - mobile, B - BGP
       D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter area
       N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type 2
       E1 - OSPF external type 1, E2 - OSPF external type 2
       i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1, L2 - IS-IS level-2
       ia - IS-IS inter area, * - candidate default, U - per-user static route
       o - ODR, P - periodic downloaded static route

Gateway of last resort is not set

    10.0.0.0/24 is subnetted, 3 subnets
B       10.3.3.0 [20/0] via 192.168.1.6, 00:04:51
B       10.2.2.0 [20/0] via 192.168.1.6, 00:05:22
C       10.1.1.0 is directly connected, Loopback0
    192.168.1.0/30 is subnetted, 1 subnets
C       192.168.1.4 is directly connected, Serial1/0

```

b. Ping again, this time as an extended ping, sourcing from the Loopback0 interface address. **ping 10.3.3.1 source 10.1.1.1 or ping 10.3.3.1 source Lo0**

```

Andheri#ping 10.3.3.1 source 10.1.1.1
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 10.3.3.1, timeout is 2 seconds:
Packet sent with a source address of 10.1.1.1
!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 36/52/64 ms
Andheri#
Andheri#show ip bgp
BGP table version is 4, local router ID is 10.1.1.1
Status codes: s suppressed, d damped, h history, * valid, > best, i - in
               r RIB-failure, S Stale
Origin codes: i - IGP, e - EGP, ? - incomplete

   Network          Next Hop          Metric LocPrf Weight Path
*> 10.1.1.0/24      0.0.0.0              0         32768 i
*> 10.2.2.0/24      192.168.1.6          0           300 i
*> 10.3.3.0/24      192.168.1.6          0           300 65000 i

```

c. Now check the BGP table on Andheri. The AS_PATH to the 10.3.3.0 network should be AS 300. It no longer has the private AS in the path.

Andheri# show ip bgp

```

Andheri#ping 10.3.3.1 source 10.1.1.1
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 10.3.3.1, timeout is 2 seconds:
Packet sent with a source address of 10.1.1.1
!!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 20/60/90
Andheri#show ip bgp
BGP table version is 5, local router ID is 10.1.1.1
Status codes: s suppressed, d damped, h history, * valid, > best, i
               r RIB-failure, S Stale
Origin codes: i - IGP, e - EGP, ? - incomplete

   Network          Next Hop          Metric LocPrf Weight Path
*> 10.1.1.0/24      0.0.0.0              0         32768 i
*> 10.2.2.0/24      192.168.1.6          0           0 300 i
*> 10.3.3.0/24      192.168.1.6          0           0 300 i

```

Step 5 : Use the AS_PATH attribute to filter routes.

- a. Configure a special kind of access list to match BGP routes with an AS_PATH attribute that both begins and ends with the number 100. Enter the following commands on Bandra.

```
Bandra(config)# ip as-path access-list 1 deny ^100$
```

```
Bandra(config)# ip as-path access-list 1 permit .*
```

```

Bandra#
Bandra#conf t
Enter configuration commands, one per line. End with CNTL/Z.
Bandra(config)#ip as-path access-list 1 deny ^100$
Bandra(config)#ip as-path access-list 1 permit .*

```

- b. Apply the configured access list using the neighbor command with the filter-list option.

```
Bandra(config)# router bgp 300
```

```
Bandra (config-router)# neighbor 192.168.1.5 remove-private-as
```

```

Bandra(config)#router bgp 300
Bandra(config-router)#neighbor 172.24.1.18 filter-list 1 out
Bandra(config-router)#exit

```

- c. Use the clear ip bgp * command to reset the routing information. Wait several seconds and then check the routing table for BANDRA. The route to 10.1.1.0 should be in the routing table.

```
Andheri# show ip route
```

```

Bandra#show ip route
Codes: C - connected, S - static, R - RIP, M - mobile, B - BGP
       D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter area
       N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type 2
       E1 - OSPF external type 1, E2 - OSPF external type 2
       i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1, L2 - IS-IS level-2
       ia - IS-IS inter area, * - candidate default, U - per-user static route
       o - ODR, P - periodic downloaded static route

Gateway of last resort is not set

    172.24.0.0/30 is subnetted, 1 subnets
C       172.24.1.16 is directly connected, Serial1/1
    10.0.0.0/24 is subnetted, 3 subnets
B       10.3.3.0 [20/0] via 172.24.1.18, 00:13:19
C       10.2.2.0 is directly connected, Loopback0
B       10.1.1.0 [20/0] via 192.168.1.5, 00:13:20
    192.168.1.0/30 is subnetted, 1 subnets
C       192.168.1.4 is directly connected, Serial1/0

```

d. Return to BANDRA and verify that the filter is working as intended.

```
Bandra# show ip bgp regexp ^100$
```

```

Bandra#show ip bgp regexp ^100$
BGP table version is 4, local router ID is 10.2.2.1
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
               r RIB-failure, S Stale
Origin codes: i - IGP, e - EGP, ? - incomplete

   Network          Next Hop           Metric LocPrf Weight Path
*> 10.1.1.0/24      192.168.1.5              0             0 100 i

```

e. Run the following Tcl script on all routers to verify whether there is connectivity. All pings from BANDRA should be successful. Andheri should not be able to ping the Churchgate loopback 10.3.3.1 or the WAN link 172.24.1.16/30. Churchgate should not be able to ping the Andheri loopback 10.1.1.1 or the WAN link 192.168.1.4/30.

```
Bandra#tclsh
```

```

Bandra(tcl)#foreach address {
+>10.1.1.1
+>10.2.2.1
+>10.3.3.1
+>192.168.1.5
+>192.168.1.6
+>172.24.1.17
+>172.24.1.18
+>} { ping $address }

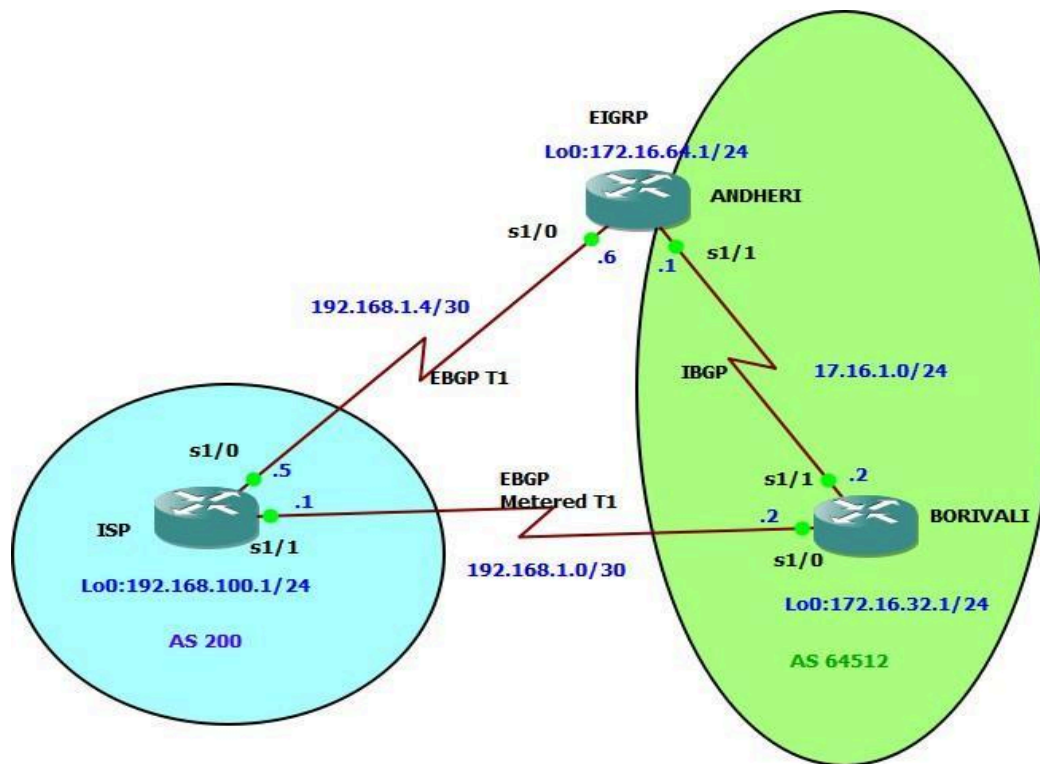
```

```
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 10.1.1.1, timeout is 2 seconds:
!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 28/40/64 ms
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 10.2.2.1, timeout is 2 seconds:
!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 4/4/4 ms
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 10.3.3.1, timeout is 2 seconds:
!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 16/32/48 ms
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 192.168.1.5, timeout is 2 seconds:
!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 16/28/40 ms
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 192.168.1.6, timeout is 2 seconds:
!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 48/59/64 ms
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 172.24.1.17, timeout is 2 seconds:
!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 48/56/68 ms
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 172.24.1.18, timeout is 2 seconds:
!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 16/31/48 ms
```


Practical No - 3

Aim: Configuring IBGP and EBGP Sessions, Local Preference, and MED

Topolgy :



Objectives

- For IBGP peers to correctly exchange routing information, use the **next-hop-self** command with the **Local-Preference** and **MED** attributes.
- Ensure that the flat-rate, unlimited-use T1 link is used for sending and receiving data to and from the AS 200 on ISP and that the metered T1 only be used in the event that the primary T1 link has failed.

Step 1: Configure interface addresses.

Router R1 (hostname ISP)

```
ISP(config)# interface Loopback0
```

```
ISP(config-if)# ip address 192.168.100.1 255.255.255.0
```

```
ISP(config-if)# exit
```

```
ISP(config)# interface Serial0/0/0
```

ISP(config-if)# **ip address 192.168.1.5 255.255.255.252**

ISP(config-if)# **no shutdown**

ISP(config-if)# **exit**

ISP(config)# **interface Serial0/0/1**

ISP(config-if)# **ip address 192.168.1.1 255.255.255.252**

ISP(config-if)# **no shutdown**

ISP(config-if)# **end**

```
ISP(config-if)#interface Loopback0
ISP(config-if)#ip address 192.168.100.1 255.255.255.0
ISP(config-if)#exit
ISP(config)#
ISP(config)#int s1/0
ISP(config-if)#ip address 192.168.1.5 255.255.255.252
ISP(config-if)#no shutdown
ISP(config-if)#exit
ISP(config)#
*May 18 17:42:51.491: %LINK-3-UPDOWN: Interface Serial11
ISP(config)#
*May 18 17:42:52.495: %LINEPROTO-5-UPDOWN: Line protoco
ISP(config)#
ISP(config)#int s1/1
ISP(config-if)#ip address 192.168.1.1 255.255.255.252
ISP(config-if)#no shutdown
```

Router R2 (hostname SanJose1)

SanJose1(config)# **interface Loopback0**

SanJose1(config-if)# **ip address 172.16.64.1 255.255.255.0**

SanJose1(config-if)# **exit**

SanJose1(config)# **interface Serial0/0/0**

SanJose1(config-if)# **ip address 192.168.1.6 255.255.255.252**

SanJose1(config-if)# **no shutdown**

SanJose1(config-if)# **exit**

SanJose1(config)# **interface Serial0/0/1**

SanJose1(config-if)# **ip address 172.16.1.1 255.255.255.0**

SanJose1(config-if)# **no shutdown**

SanJose1(config-if)# **end**

```
ANDHERI(config)#interface Loopback0
ANDHERI(config-if)#
*May 18 17:42:40.167: %LINEPROTO-5-UPDOWN: Line protocol on Interface Loopback0, line protocol is up
ANDHERI(config-if)#ip address 172.16.64.1 255.255.255.0
ANDHERI(config-if)#exit
ANDHERI(config)#
ANDHERI(config)#int s1/0
ANDHERI(config-if)#ip address 192.168.1.6 255.255.255.252
ANDHERI(config-if)#no shutdown
ANDHERI(config-if)#exit
ANDHERI(config)#
*May 18 17:43:11.899: %LINK-3-UPDOWN: Interface Serial1/0/0, line protocol is down
ANDHERI(config)#
*May 18 17:43:12.903: %LINEPROTO-5-UPDOWN: Line protocol on Interface Serial1/0/0, line protocol is down
ANDHERI(config)#
ANDHERI(config)#int s1/1
ANDHERI(config-if)#ip address 172.16.1.1 255.255.255.0
ANDHERI(config-if)#no shutdown
```

Router R3 (hostname SanJose2)

```
SanJose2(config)# interface Loopback0

SanJose2(config-if)# ip address 172.16.32.1 255.255.255.0

SanJose2(config-if)# exit

SanJose2(config)# interface Serial0/0/0

SanJose2(config-if)# ip address 192.168.1.2 255.255.255.252

SanJose2(config-if)# no shutdown

SanJose2(config-if)# exit

SanJose2(config)# interface Serial0/0/1

SanJose2(config-if)# ip address 172.16.1.2 255.255.255.0

SanJose2(config-if)# no shutdown

SanJose2(config-if)# end
```

```

BORIVALI(config)#interface Loopback0
BORIVALI(config-if)#
*May 18 17:43:25.783: %LINEPROTO-5-UPDOWN: Line protocol o
BORIVALI(config-if)#ip address 172.16.32.1 255.255.255.0
BORIVALI(config-if)#exit
BORIVALI(config)#
BORIVALI(config)#int s1/0
BORIVALI(config-if)#ip address 192.168.1.2 255.255.255.252
BORIVALI(config-if)#no shutdown
BORIVALI(config-if)#exit
BORIVALI(config)#
*May 18 17:43:54.311: %LINK-3-UPDOWN: Interface Serial1/0,
BORIVALI(config)#
BORIVALI(config)#
*May 18 17:43:55.315: %LINEPROTO-5-UPDOWN: Line protocol o
BORIVALI(config)#int s1/1
BORIVALI(config-if)#ip address 172.16.1.2 255.255.255.0
BORIVALI(config-if)#no shutdown

```

Step 2: Configure EIGRP.

Configure EIGRP between the SanJose1 and SanJose2 routers. (Note: If using an IOS prior to 15.0, use the no auto-summary router configuration command to disable automatic summarization. This command is the default beginning with IOS 15.)

```
SanJose1(config)# router eigrp 1
```

```
SanJose1(config-router)# network 172.16.0.0
```

```

ANDHERI(config)# router eigrp 1
ANDHERI(config-router)#network 172.16.0.0

```

```
SanJose2(config)# router eigrp 1
```

```
SanJose2(config-router)# network 172.16.0.0
```

```

BORIVALI(config)#router eigrp 1
BORIVALI(config-router)#network 172.16.0.0

```

Step 3: Configure IBGP and verify BGP neighbors.

- Configure IBGP between the SanJose1 and SanJose2 routers. On the SanJose1 router, enter the following configuration.

```
SanJose1(config)# router bgp 64512
```

```
SanJose1(config-router)# neighbor 172.16.32.1 remote-as 64512
```

```
SanJose1(config-router)# neighbor 172.16.32.1 update-source lo0
```

```

ANDHERI(config)#
ANDHERI(config)#router bgp 64512
ANDHERI(config-router)#neighbor 172.16.32.1 remote-as 64512
ANDHERI(config-router)#neighbor 172.16.32.1 update-source lo0

```

If multiple pathways to the BGP neighbor exist, the router can use multiple IP interfaces to communicate with the neighbor. The source IP address therefore depends on the outgoing interface. The **update-source**

lo0 command instructs the router to use the IP address of the interface Loopback0 as the source IP address for all BGP messages sent to that neighbor.

- b. Complete the IBGP configuration on SanJose2 using the following commands.

```
SanJose2(config)# router bgp 64512
```

```
SanJose2(config-router)# neighbor 172.16.64.1 remote-as 64512
```

```
SanJose2(config-router)# neighbor 172.16.64.1 update-source lo0
```

```
BORIVALI(config)#router bgp 64512
BORIVALI(config-router)#neighbor 172.16.64.1 remote-as 64512
BORIVALI(config-router)#neighbor 172.16.64.1 update-source lo0
```

- c. Verify that SanJose1 and SanJose2 become BGP neighbors by issuing the **show ip bgp neighbors** command on SanJose1. View the following partial output. If the BGP state is not established, troubleshoot the connection.

```
SanJose2# show ip bgp neighbors
```

```
BORIVALI#show ip bgp neighbors
BGP neighbor is 172.16.64.1, remote AS 64512,
  BGP version 4, remote router ID 172.16.64.1
  BGP state = Established, up for 00:00:47
  Last read 00:00:47, last write 00:00:47, hold time is 180
  Neighbor capabilities:
    Route refresh: advertised and received(new
    New ASN Capability: advertised and receive
    Address family IPv4 Unicast: advertised and
  Message statistics:
    InQ depth is 0
    OutQ depth is 0

              Sent      Rcvd
Opens:             1         1
Notifications:     0         0
Updates:           0         0
```

Step 4: Configure EBGp and verify BGP neighbors.

- d. Configure ISP to run EBGp with SanJose1 and SanJose2. Enter the following commands on ISP.

```
ISP(config)# router bgp 200
```

```
ISP(config-router)# neighbor 192.168.1.6 remote-as 64512
```

```
ISP(config-router)# neighbor 192.168.1.2 remote-as 64512
```

```
ISP(config-router)# network 192.168.100.0
```

```
ISP(config)#router bgp 200
ISP(config-router)#neighbor 192.168.1.6 remote-as 64512
ISP(config-router)#neighbor 192.168.1.2 remote-as 64512
ISP(config-router)#network 192.168.100.0
```

- e. Configure a discard static route for the 172.16.0.0/16 network. Any packets that do not have a more specific match (longer match) for a 172.16.0.0 subnet will be dropped instead of sent to the ISP. Later in this lab we will configure a default route to the ISP.

```
SanJose1(config)# ip route 172.16.0.0 255.255.0.0 null0
```

```
ANDHERI(config)#ip route 172.16.0.0 255.255.0.0 null0
ANDHERI(config)#
```

- f. Configure SanJose1 as an EBGp peer to ISP.

```
SanJose1(config)# router bgp 64512
```

```
SanJose1(config-router)# neighbor 192.168.1.5 remote-as 200
```

```
SanJose1(config-router)# network 172.16.0.0
```

```
ANDHERI(config)#router bgp 64512
ANDHERI(config-router)#neighbor 192.168.1.5 remote-as 200
ANDHERI(config-router)#network 172.16.0.0
ANDHERI(config-router)#exit
```

- g. Use the **show ip bgp neighbors** command to verify that SanJose1 and ISP have reached the established state. Troubleshoot if necessary.

```
SanJose1# show ip bgp neighbors
```

```
ANDHERI#show ip bgp neighbors
BGP neighbor is 172.16.32.1, remote AS 64512, internal link
  BGP version 4, remote router ID 172.16.32.1
  BGP state = Established, up for 00:02:49
  Last read 00:00:56, last write 00:00:21, hold time is 180
  Neighbor capabilities:
    Route refresh: advertised and received(new)
    New ASN Capability: advertised and received
    Address family IPv4 Unicast: advertised and received
  Message statistics:
    InQ depth is 0
    OutQ depth is 0
```

Configure a discard static route for 172.16.0.0/16 on SanJose2 and as an EBGp peer to ISP.

```
SanJose2(config)# ip route 172.16.0.0 255.255.0.0 null0
```

```
SanJose2(config)# router bgp 64512
```

```
SanJose2(config-router)# neighbor 192.168.1.1 remote-as 200
```

```
SanJose2(config-router)# network 172.16.0.0
```

```

BORIVALI(config)#ip route 172.16.0.0 255.255.0.0 null0
BORIVALI(config)#router bgp 64512
BORIVALI(config-router)#neighbor 192.168.1.1 remote-as 200
BORIVALI(config-router)#
*May 18 18:00:01.031: %BGP-5-ADJCHANGE: neighbor 192.168.1.1 Up
BORIVALI(config-router)#network 172.16.0.0

```

Step 5: View BGP summary output.

In Step 4, the **show ip bgp neighbors** command was used to verify that SanJose1 and ISP had reached the established state. A useful alternative command is **show ip bgp summary**. The output should be similar to the following.

SanJose2# **show ip bgp summary**

```

BORIVALI# show ip bgp summary
BGP router identifier 172.16.32.1, local AS number 64512
BGP table version is 5, main routing table version 5
2 network entries using 264 bytes of memory
4 path entries using 208 bytes of memory
5/2 BGP path/bestpath attribute entries using 840 bytes of memory
1 BGP AS-PATH entries using 24 bytes of memory
0 BGP route-map cache entries using 0 bytes of memory
0 BGP filter-list cache entries using 0 bytes of memory
Bitfield cache entries: current 2 (at peak 2) using 64 bytes of memory
BGP using 1400 total bytes of memory
BGP activity 2/0 prefixes, 4/0 paths, scan interval 60 secs

Neighbor        V    AS MsgRcvd MsgSent   TblVer   InQ  OutQ Up/Down   State/Pfx
172.16.64.1      4   64512      7      8        5     0    0 00:04:18    2
192.168.1.1      4    200      5      4        3     0    0 00:00:26    1

```

Step 6: Verify which path the traffic takes.

f. Clear the IP BGP conversation with the **clear ip bgp *** command on ISP. Wait for the conversations to reestablish with each SanJose router.

ISP# **clear ip bgp ***

```

ISP#clear ip bgp *
ISP#

```

g. Test whether ISP can ping the loopback 0 address of 172.16.64.1 on SanJose1 and the serial link between SanJose1 and SanJose2, 172.16.1.1.

ISP# **ping 172.16.64.1**

```

ISP#ping 172.16.64.1

Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 172.16.64.1, timeout is 2 seconds:

*May 18 18:02:42.575: %BGP-5-ADJCHANGE: neighbor 192.168.1.6 Up .....
Success rate is 0 percent (0/5)

```

ISP# **ping 172.16.1.1**

```
ISP#ping 172.16.1.1

Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 172.16.1.1, timeout is 2 seconds:
.....
Success rate is 0 percent (0/5)
```

- h. Now ping from ISP to the loopback 0 address of 172.16.32.1 on SanJose2 and the serial link between SanJose1 and SanJose2, 172.16.1.2.

ISP# ping 172.16.32.1

```
ISP# ping 172.16.32.1

Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 172.16.32.1, timeout is 2 seconds:
!!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 12/31/48 ms
```

ISP# ping 172.16.1.2

```
ISP#ping 172.16.1.2

Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 172.16.1.2, timeout is 2 seconds:
!!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 16/24/32 ms
```

- i. Issue the **show ip bgp** command on ISP to verify BGP routes and metrics.

ISP# show ip bgp

```
ISP#show ip bgp
BGP table version is 3, local router ID is 192.168.100.1
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
               r RIB-failure, S Stale
Origin codes: i - IGP, e - EGP, ? - incomplete

   Network        Next Hop        Metric LocPrf Weight Path
* 172.16.0.0      192.168.1.6          0           0 64512 i
*>                192.168.1.2          0           0 64512 i
*> 192.168.100.0  0.0.0.0              0          32768 i
```

- i. At this point, the ISP router should be able to get to each network connected to SanJose1 and SanJose2 from the loopback address 192.168.100.1. Use the extended **ping** command and specify the source address of ISP Lo0 to test.

ISP# ping 172.16.1.1 source 192.168.100.1

```
ISP#ping 172.16.1.1 source 192.168.100.1

Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 172.16.1.1, timeout is 2 seconds:
Packet sent with a source address of 192.168.100.1
!!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 16/38/48 ms
```

ISP# ping 172.16.32.1 source 192.168.100.1


```
ISP# ping 172.16.32.1 source 192.168.100.1
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 172.16.32.1, timeout is 2 seconds:
Packet sent with a source address of 192.168.100.1
!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 4/32/48 ms
```

ISP# ping 172.16.1.2 source 192.168.100.1

```
ISP#ping 172.16.1.2 source 192.168.100.1
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 172.16.1.2, timeout is 2 seconds:
Packet sent with a source address of 192.168.100.1
!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 16/28/48 ms
```

ISP# ping 172.16.64.1 source 192.168.100.1

```
ISP#ping 172.16.64.1 source 192.168.100.1
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 172.16.64.1, timeout is 2 seconds:
Packet sent with a source address of 192.168.100.1
!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 24/35/56 ms
```

Step 7: Configure the BGP next-hop-self feature.

- j. Issue the following commands on the ISP router.

```
ISP(config)# router bgp 200
```

```
ISP(config-router)# network 192.168.1.0 mask 255.255.255.252
```

```
ISP(config-router)# network 192.168.1.4 mask 255.255.255.252
```

```
ISP(config)#router bgp 200
ISP(config-router)#network 192.168.1.0 mask 255.255.255.252
ISP(config-router)#network 192.168.1.4 mask 255.255.255.252
ISP(config-router)#exit
```

- k. Issue the **show ip bgp** command to verify that the ISP is correctly injecting its own WAN links into BGP.

```
ISP# show ip bgp
```

```
ISP#show ip bgp
BGP table version is 5, local router ID is 192.168.100.1
Status codes: s suppressed, d damped, h history, * valid, > best, i -
               r RIB-failure, S Stale
Origin codes: i - IGP, e - EGP, ? - incomplete

   Network          Next Hop          Metric LocPrf Weight Path
* 172.16.0.0         192.168.1.6             0           0 64512 i
*>                  192.168.1.2             0           0 64512 i
*> 192.168.1.0/30    0.0.0.0                 0          32768 i
*> 192.168.1.4/30    0.0.0.0                 0          32768 i
*> 192.168.100.0     0.0.0.0                 0          32768 i
```

- l. Verify on SanJose1 and SanJose2 that the opposite WAN link is included in the routing table. The output from SanJose2 is as follows.

SanJose2# **show ip route**

```
BORIVALI#show ip route
Codes: C - connected, S - static, R - RIP, M - mobile, B - BGP
       D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter area
       N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type 2
       E1 - OSPF external type 1, E2 - OSPF external type 2
       i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1, L2 - IS-IS level-2
       ia - IS-IS inter area, * - candidate default, U - per-user static route
       o - ODR, P - periodic downloaded static route
```

- m. To better understand the **next-hop-self** command we will remove ISP advertising its two WAN links and shutdown the WAN link between ISP and SanJose2. The only possible path from SanJose2 to ISP's 192.168.100.0/24 is through SanJose1.

ISP(config)# **router bgp 200**

ISP(config-router)# **no network 192.168.1.0 mask 255.255.255.252**

ISP(config-router)# **no network 192.168.1.4 mask 255.255.255.252**

ISP(config-router)# **exit**

ISP(config)# **interface serial 0/0/1**

ISP(config-if)# **shutdown**

```
ISP(config)#router bgp 200
ISP(config-router)#no network 192.168.1.0 mask 255.255.255.252
ISP(config-router)#no network 192.168.1.4 mask 255.255.255.252
ISP(config-router)#exit
ISP(config)#int s1/1
ISP(config-if)#shutdown
```

- n. Display SanJose2's BGP table using the **show ip bgp** command and the IPv4 routing table with **show ip route**.

SanJose2# **show ip bgp**

```

BORIVALI#show ip bgp
BGP table version is 14, local router ID is 172.16.32.1
Status codes: s suppressed, d damped, h history, * valid, > best, i -
               r RIB-failure, S Stale
Origin codes: i - IGP, e - EGP, ? - incomplete

   Network          Next Hop              Metric LocPrf Weight Path
*> 172.16.0.0        0.0.0.0                  0         32768 i
* i                 172.16.64.1              0        100      0 i
* i192.168.100.0     192.168.1.5              0        100      0 200 i
*>                  192.168.1.1              0         0 200 i

```

SanJose2# **show ip route**

```

BORIVALI#show ip route
Codes: C - connected, S - static, R - RIP, M - mobile, B - BGP
       D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter area
       N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type 2
       E1 - OSPF external type 1, E2 - OSPF external type 2
       i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1, L2 - IS-IS level-2
       ia - IS-IS inter area, * - candidate default, U - per-user static route
       o - ODR, P - periodic downloaded static route

Gateway of last resort is not set

 172.16.0.0/16 is variably subnetted, 4 subnets, 2 masks
C    172.16.32.0/24 is directly connected, Loopback0
S    172.16.0.0/16 is directly connected, Null0
C    172.16.1.0/24 is directly connected, Serial1/1
D    172.16.64.0/24 [90/2297856] via 172.16.1.1, 00:11:47, Serial1/1

```

SanJose1(config)# **router bgp 64512**

SanJose1(config-router)# **neighbor 172.16.32.1 next-hop-self**

```

ANDHERI(config)#router bgp 64512
ANDHERI(config-router)#neighbor 172.16.32.1 next-hop-self
ANDHERI(config-router)#exit

```

SanJose2(config)# **router bgp 64512**

SanJose2(config-router)# **neighbor 172.16.64.1 next-hop-self**

```

BORIVALI(config-router)#router bgp 64512
BORIVALI(config-router)#neighbor 172.16.64.1 next-hop-self

```

- o. Reset BGP operation on either router with the **clear ip bgp *** command.

SanJose1# **clear ip bgp ***

```

ANDHERI#clear ip bgp *

```

SanJose2# **clear ip bgp ***

```

BORIVALI#clear ip bgp *

```

- p. After the routers have returned to established BGP speakers, issue the **show ip bgp** command on SanJose2 and notice that the next hop is now SanJose1 instead of ISP.

SanJose2# **show ip bgp**

```
BORIVALI#show ip bgp
BGP table version is 1, local router ID is 172.16.32.1
Status codes: s suppressed, d damped, h history, * valid, > best, i
               r RIB-failure, S Stale
Origin codes: i - IGP, e - EGP, ? - incomplete

   Network          Next Hop          Metric LocPrf Weight Path
* i172.16.0.0        172.16.64.1              0    100      0 i
* i192.168.100.0     172.16.64.1              0    100      0 200 i
```

- q. The **show ip route** command on SanJose2 now displays the 192.168.100.0/24 network because SanJose1 is the next hop, 172.16.64.1, which is reachable from SanJose2.

SanJose2# **show ip route**

```
BORIVALI#show ip route
Codes: C - connected, S - static, R - RIP, M - mobile, B - BGP
       D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter area
       N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type 2
       E1 - OSPF external type 1, E2 - OSPF external type 2
       i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1, L2 - IS-IS level
       ia - IS-IS inter area, * - candidate default, U - per-user static r
       o - ODR, P - periodic downloaded static route

Gateway of last resort is not set
```

- r. Before configuring the next BGP attribute, restore the WAN link between ISP and SanJose3. This will change the BGP table and routing table on both routers. For example, SanJose2's routing table shows 192.168.100.0/24 will now have a better path through ISP.

ISP(config)# **interface serial 0/0/1**

ISP(config-if)# **no shutdown**

```
ISP(config)#int s1/1
ISP(config-if)#no shutdown
```

SanJose2# **show ip route**

```

BORIVALI#show ip route
Codes: C - connected, S - static, R - RIP, M - mobile, B - BGP
        D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter area
        N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type 2
        E1 - OSPF external type 1, E2 - OSPF external type 2
        i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1, L2 - IS-IS level-2
        ia - IS-IS inter area, * - candidate default, U - per default
        o - ODR, P - periodic downloaded static route

Gateway of last resort is not set

    172.16.0.0/16 is variably subnetted, 4 subnets, 2 masks
C       172.16.32.0/24 is directly connected, Loopback0
S       172.16.0.0/16 is directly connected, Null0
C       172.16.1.0/24 is directly connected, Serial1/1
D       172.16.64.0/24 [90/2297856] via 172.16.1.1, 00:15:05, Serial1/0
    192.168.1.0/30 is subnetted, 1 subnets
C       192.168.1.0 is directly connected, Serial1/0
B       192.168.100.0/24 [20/0] via 192.168.1.1, 00:00:18

```

Step 8: Set BGP local preference.

- s. Because the local preference value is shared between IBGP neighbors, configure a simple route map that references the local preference value on SanJose1 and SanJose2. This policy adjusts outbound traffic to prefer the link off the SanJose1 router instead of the metered T1 off SanJose2.

```
SanJose1(config)# route-map PRIMARY_T1_IN permit 10
```

```
SanJose1(config-route-map)# set local-preference 150
```

```
SanJose1(config-route-map)# exit
```

```
SanJose1(config)# router bgp 64512
```

```
SanJose1(config-router)# neighbor 192.168.1.5 route-map PRIMARY_T1_IN in
```

```

ANDHERI(config)#route-map PRIMARY_T1_IN permit 10
ANDHERI(config-route-map)#set local-preference 150
ANDHERI(config-route-map)#exit
ANDHERI(config)#router bgp 64512
ANDHERI(config-router)#neighbor 192.168.1.5 route-map PRIMARY_T1_IN in

```

```
SanJose2(config)# route-map SECONDARY_T1_IN permit 10
```

```
SanJose2(config-route-map)# set local-preference 125
```

```
SanJose1(config-route-map)# exit
```

```
SanJose2(config)# router bgp 64512
```

```
SanJose2(config-router)# neighbor 192.168.1.1 route-map SECONDARY_T1_IN in
```

```

BORIVALI(config)#route-map SECONDARY_T1_IN permit 10
BORIVALI(config-route-map)#set local-preference 125
BORIVALI(config-route-map)#exit
BORIVALI(config)#router bgp 64512
BORIVALI(config-router)#neighbor 192.168.1.1 route-map SECONDARY_T1_IN in

```

- t. Use the **clear ip bgp * soft** command after configuring this new policy. When the conversations have been reestablished, issue the **show ip bgp** command on SanJose1 and SanJose2.

SanJose1# **clear ip bgp * soft**

```
ANDHERI#clear ip bgp * soft
```

SanJose2# **clear ip bgp * soft**

```
BORIVALI#clear ip bgp * soft
```

SanJose1# **show ip bgp**

```
ANDHERI#show ip bgp
BGP table version is 6, local router ID is 172.16.64.1
Status codes: s suppressed, d damped, h history, * valid, > best, i
               r RIB-failure, S Stale
Origin codes: i - IGP, e - EGP, ? - incomplete

   Network          Next Hop          Metric LocPrf Weight Path
* i172.16.0.0        172.16.32.1              0    100      0 i
*>                   0.0.0.0                  0          32768 i
*> 192.168.100.0     192.168.1.5              0    150      0 200 i
```

SanJose2# **show ip bgp**

```
BORIVALI#show ip bgp
BGP table version is 5, local router ID is 172.16.32.1
Status codes: s suppressed, d damped, h history, * valid, > best, i
               r RIB-failure, S Stale
Origin codes: i - IGP, e - EGP, ? - incomplete

   Network          Next Hop          Metric LocPrf Weight Path
*> 172.16.0.0        0.0.0.0              0          32768 i
* i                  172.16.64.1          0    100      0 i
* 192.168.100.0     192.168.1.1          0          0 200 i
*>i                  172.16.64.1          0    150      0 200 i
```

Step 9: Set BGP MED.

- u. In the previous step we saw that SanJose1 and SanJose2 will route traffic for 192.168.100.0/24 using the link between SanJose1 and ISP. Examine what the return path ISP takes to reach AS 64512. Notice that the return path is different from the original path. This is known as asymmetric routing and is not necessarily an unwanted trait.

ISP# **show ip bgp**

```

ISP#show ip bgp
BGP table version is 11, local router ID is 192.168.100.1
Status codes: s suppressed, d damped, h history, * valid, > best, i -
               r RIB-failure, S Stale
Origin codes: i - IGP, e - EGP, ? - incomplete

   Network          Next Hop          Metric LocPrf Weight Path
* 172.16.0.0        192.168.1.2            0           0 64512 i
*>                  192.168.1.6            0           0 64512 i
*> 192.168.100.0    0.0.0.0                0          32768 i

```

ISP# show ip route

```

ISP#show ip route
Codes: C - connected, S - static, R - RIP, M - mobile, B
       D - EIGRP, EX - EIGRP external, O - OSPF, IA - OS
       N1 - OSPF NSSA external type 1, N2 - OSPF NSSA ex
       E1 - OSPF external type 1, E2 - OSPF external typ
       i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1
       ia - IS-IS inter area, * - candidate default, U -
       o - ODR, P - periodic downloaded static route

Gateway of last resort is not set

B    172.16.0.0/16 [20/0] via 192.168.1.6, 00:04:56
     192.168.1.0/30 is subnetted, 2 subnets
C      192.168.1.0 is directly connected, Serial1/1
C      192.168.1.4 is directly connected, Serial1/0
C    192.168.100.0/24 is directly connected, Loopback0

```

- a. Use an extended **ping** command to verify this situation. Specify the **record** option and compare your output to the following. Notice the return path using the exit interface 192.168.1.1 to SanJose2

```

BORIVALI#ping
Protocol [ip]:
Target IP address: 192.168.100.1
Repeat count [5]:
Datagram size [100]:
Timeout in seconds [2]:
Extended commands [n]: y
Source address or interface: 172.16.32.1
Type of service [0]:
Set DF bit in IP header? [no]:
Validate reply data? [no]:
Data pattern [0xABCD]:
Loose, Strict, Record, Timestamp, Verbose[none]: record
Number of hops [ 9 ]:
Loose, Strict, Record, Timestamp, Verbose[RV]:
Sweep range of sizes [n]:
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 192.168.100.1, timeout is 2 seconds:
Packet sent with a source address of 172.16.32.1
Packet has IP options: Total option bytes= 39, padded length=40
Record route: <*>
(0.0.0.0)
(0.0.0.0)
(0.0.0.0)
(0.0.0.0)
(0.0.0.0)
(0.0.0.0)
(0.0.0.0)
(0.0.0.0)
(0.0.0.0)

```

```

Reply to request 4 (8 ms). Received packet has options
Total option bytes= 40, padded length=40
Record route:
(172.16.1.2)
(192.168.1.6)
(192.168.100.1)
(192.168.1.5)
(172.16.1.1)
(172.16.32.1) <*>
(0.0.0.0)
(0.0.0.0)
(0.0.0.0)
End of list

Success rate is 100 percent (5/5), round-trip min/avg/max = 8/47/68 ms

```

If you are unfamiliar with the **record** option, the important thing to note is that each IP address in brackets is an outgoing interface. The output can be interpreted as follows:

1. A ping that is sourced from 172.16.32.1 exits SanJose2 through s0/0/1, 172.16.1.2. It then arrives at the s0/0/1 interface for SanJose1.
2. SanJose1 S0/0/0, 192.168.1.6, routes the packet out to arrive at the S0/0/0 interface of ISP.
3. The target of 192.168.100.1 is reached: 192.168.100.1.
4. The packet is next forwarded out the S0/0/1, 192.168.1.1 interface for ISP and arrives at the S0/0/0 interface for SanJose2.
5. SanJose2 then forwards the packet out the last interface, loopback 0, 172.16.32.1.

Although the unlimited use of the T1 from SanJose1 is preferred here, ISP currently takes the link from SanJose2 for all return traffic.

- b. Create a new policy to force the ISP router to return all traffic via SanJose1. Create a second route map utilizing the MED (metric) that is shared between EBGp neighbors.

```
SanJose1(config)#route-map PRIMARY_T1_MED_OUT permit 10
```

```
SanJose1(config-route-map)#set Metric 50
```

```
SanJose1(config-route-map)#exit
```

```
SanJose1(config)#router bgp 64512
```

```
SanJose1(config-router)#neighbor 192.168.1.5 route-map PRIMARY_T1_MED_OUT out
```

```
ANDHERI(config)#route-map PRIMARY_T1_MED_OUT permit 10
ANDHERI(config-route-map)#set Metric 50
ANDHERI(config-route-map)#exit
ANDHERI(config)#router bgp 64512
ANDHERI(config-router)#neighbor 192.168.1.5 route-map PRIMARY_T1_MED_OUT out
```

```
SanJose2(config)#route-map SECONDARY_T1_MED_OUT permit 10
```

```
SanJose2(config-route-map)#set Metric 75
```

```
SanJose2(config-route-map)#exit
```

```
SanJose2(config)#router bgp 64512
```

```
SanJose2(config-router)#neighbor 192.168.1.1 route-map SECONDARY_T1_MED_OUT out
```

```
BORIVALI(config)#route-map SECONDARY_T1_MED_OUT permit 10
BORIVALI(config-route-map)#set Metric 75
BORIVALI(config-route-map)#exit
BORIVALI(config)#router bgp 64512
BORIVALI(config-router)#neighbor 192.168.1.1 route-map SECONDARY_T1_MED_OUT out
```

- v. Use the **clear ip bgp * soft** command after issuing this new policy. Issuing the **show ip bgp** command as follows on SanJose1 or SanJose2 does not indicate anything about this newly defined policy.

```
SanJose1# clear ip bgp * soft
```

```
ANDHERI#clear ip bgp * soft
```

```
SanJose2# clear ip bgp * soft
```

```
BORIVALI#clear ip bgp * soft
```

```
SanJose1# show ip bgp
```

```

ANDHERI#show ip bgp
BGP table version is 6, local router ID is 172.16.64.1
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
               r RIB-failure, S Stale
Origin codes: i - IGP, e - EGP, ? - incomplete

   Network          Next Hop          Metric LocPrf Weight Path
* i172.16.0.0        172.16.32.1              0    100      0 i
*>                  0.0.0.0                  0          32768 i
*> 192.168.100.0     192.168.1.5              0    150      0 200 i

```

SanJose2# show ip bgp

```

BORIVALI#show ip bgp
BGP table version is 5, local router ID is 172.16.32.1
Status codes: s suppressed, d damped, h history, * valid, > best, i
               r RIB-failure, S Stale
Origin codes: i - IGP, e - EGP, ? - incomplete

   Network          Next Hop          Metric LocPrf Weight Path
*> 172.16.0.0        0.0.0.0              0          32768 i
* i                  172.16.64.1          0    100      0 i
* 192.168.100.0     192.168.1.1          0    125      0 200 i
*>i                  172.16.64.1          0    150      0 200 i

```

Reissue an extended **ping** command with the **record** command. Notice the change in return path using the exit interface 192.168.1.5 to SanJose1.

```

BORIVALI#ping
Protocol [ip]:
Target IP address: 192.168.100.1
Repeat count [5]:
Datagram size [100]:
Timeout in seconds [2]:
Extended commands [n]: y
Source address or interface: 172.16.32.1
Type of service [0]:
Set DF bit in IP header? [no]:
Validate reply data? [no]:
Data pattern [0xABCD]:
Loose, Strict, Record, Timestamp, Verbose[none]: record
Number of hops [ 9 ]:
Loose, Strict, Record, Timestamp, Verbose[RV]:
Sweep range of sizes [n]:
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 192.168.100.1, timeout is 2 sec
Packet sent with a source address of 172.16.32.1
Packet has IP options: Total option bytes= 39, padded length=40

```

```

Reply to request 3 (60 ms). Received packet has options
Total option bytes= 40, padded length=40
Record route:
(172.16.1.2)
(192.168.1.6)
(192.168.100.1)
(192.168.1.5)
(172.16.1.1)
(172.16.32.1) <*>
(0.0.0.0)
(0.0.0.0)
(0.0.0.0)
End of list

Reply to request 4 (52 ms). Received packet has options
Total option bytes= 40, padded length=40
Record route:
(172.16.1.2)
(192.168.1.6)
(192.168.100.1)
(192.168.1.5)
(172.16.1.1)
(172.16.32.1) <*>
(0.0.0.0)
(0.0.0.0)
(0.0.0.0)
End of list

Success rate is 100 percent (5/5), round-trip min/avg/max = 40/52/60 ms

```

ISP# show ip bgp

```

ISP#show ip bgp
BGP table version is 11, local router ID is 192.168.100.1
Status codes: s suppressed, d damped, h history, * valid, > best, i -
               r RIB-failure, S Stale
Origin codes: i - IGP, e - EGP, ? - incomplete

   Network          Next Hop          Metric LocPrf Weight Path
*  172.16.0.0        192.168.1.2            0           0 64512 i
*> 172.16.0.0        192.168.1.6            0           0 64512 i
*> 192.168.100.0     0.0.0.0                0           32768 i

```

Step 10: Establish a default route.

The final step is to establish a default route that uses a policy statement that adjusts to changes in the network.

- Configure ISP to inject a default route to both SanJose1 and SanJose2 using BGP using the **default-originate** command. This command does not require the presence of 0.0.0.0 in the ISP router. Configure the 10.0.0.0/8 network which will not be advertised using BGP. This network will be used to test the default route on SanJose1 and SanJose2.

ISP(config)# **router bgp 200**

ISP(config-router)# **neighbor 192.168.1.6 default-originate**

ISP(config-router)# **neighbor 192.168.1.2 default-originate**

```
ISP(config-router)# exit
```

```
ISP(config)# interface loopback 10
```

```
ISP(config-if)# ip address 10.0.0.1 255.255.255.0
```

```
ISP(config)#router bgp 200
ISP(config-router)#neighbor 192.168.1.6 default-originate
ISP(config-router)#neighbor 192.168.1.2 default-originate
ISP(config-router)#exit
```

```
ISP(config-if)#ip address 10.0.0.1 255.255.255.0
ISP(config-if)#exit
```

- b. Verify that both routers have received the default route by examining the routing tables on SanJose1 and SanJose2. Notice that both routers prefer the route between SanJose1 and ISP.

```
SanJose1# show ip route
```

```
ANDHERI#show ip route
Codes: C - connected, S - static, R - RIP, M - mobile, B
       D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF
       N1 - OSPF NSSA external type 1, N2 - OSPF NSSA ext
       E1 - OSPF external type 1, E2 - OSPF external type
       i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1,
       ia - IS-IS inter area, * - candidate default, U -
       o - ODR, P - periodic downloaded static route

Gateway of last resort is 172.16.32.1 to network 0.0.0.0

    172.16.0.0/16 is variably subnetted, 4 subnets, 2 masks
D       172.16.32.0/24 [90/2297856] via 172.16.1.2, 02:14:00
S       172.16.0.0/16 is directly connected, Null0
C       172.16.1.0/24 is directly connected, Serial1/1
C       172.16.64.0/24 is directly connected, Loopback0
B       192.168.100.0/24 [200/0] via 172.16.32.1, 01:44:43
B*     0.0.0.0/0 [200/0] via 172.16.32.1, 01:44:43
```

```
SanJose2# show ip route
```

```
BORIVALI#show ip route
Codes: C - connected, S - static, R - RIP, M - mobile, B
       D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF
       N1 - OSPF NSSA external type 1, N2 - OSPF NSSA ext
       E1 - OSPF external type 1, E2 - OSPF external type
       i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1,
       ia - IS-IS inter area, * - candidate default, U -
       o - ODR, P - periodic downloaded static route

Gateway of last resort is 192.168.1.1 to network 0.0.0.0

    172.16.0.0/16 is variably subnetted, 4 subnets, 2 masks
C       172.16.32.0/24 is directly connected, Loopback0
S       172.16.0.0/16 is directly connected, Null0
C       172.16.1.0/24 is directly connected, Serial1/1
D       172.16.64.0/24 [90/2297856] via 172.16.1.1, 02:14:00
       192.168.1.0/30 is subnetted, 1 subnets
C       192.168.1.0 is directly connected, Serial1/0
B       192.168.100.0/24 [20/0] via 192.168.1.1, 01:44:59
B*     0.0.0.0/0 [20/0] via 192.168.1.1, 01:45:00
```

- c. The preferred default route is by way of SanJose1 because of the higher local preference attribute configured on SanJose1 earlier.

SanJose2# **show ip bgp**

```
BORIVALI# show ip bgp
BGP table version is 9, local router ID is 172.16.32.1
Status codes: s suppressed, d damped, h history, * valid, > best, i -
               r RIB-failure, S Stale
Origin codes: i - IGP, e - EGP, ? - incomplete

   Network          Next Hop          Metric LocPrf Weight Path
*> 0.0.0.0           192.168.1.1           0      125      0 200 i
*> 172.16.0.0         0.0.0.0               0              32768 i
* i                   172.16.64.1           0      100       0 i
*> 192.168.100.0      192.168.1.1           0      125      0 200 i
```

- d. Using the traceroute command verify that packets to 10.0.0.1 is using the default route through SanJose1.

SanJose2# **traceroute 10.0.0.1**

```
BORIVALI#traceroute 10.0.0.1

Type escape sequence to abort.
Tracing the route to 10.0.0.1

 0 192.168.1.1 [AS 200] 28 msec 32 msec 32 msec
 1 192.168.1.1 [AS 200] !H !H !H
```

- e. Next, test how BGP adapts to using a different default route when the path between SanJose1 and ISP goes down.

ISP(config)# **interface serial 0/0/0**

ISP(config-if)# **shutdown**

```
ISP(config)#int s1/0
ISP(config-if)#shutdown
```

- f. Verify that both routers are modified their routing tables with the default route using the path between SanJose2 and ISP.

SanJose1# **show ip route**

```
ANDHERI#show ip route
Codes: C - connected, S - static, R - RIP, M - mobile, B
       D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF
       N1 - OSPF NSSA external type 1, N2 - OSPF NSSA ext
       E1 - OSPF external type 1, E2 - OSPF external type
       i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1,
       ia - IS-IS inter area, * - candidate default, U -
       o - ODR, P - periodic downloaded static route

Gateway of last resort is 172.16.32.1 to network 0.0.0.0

    172.16.0.0/16 is variably subnetted, 4 subnets, 2 ma
    D    172.16.32.0/24 [90/2297856] via 172.16.1.2, 02:15
    S    172.16.0.0/16 is directly connected, Null0
    C    172.16.1.0/24 is directly connected, Serial1/1
    C    172.16.64.0/24 is directly connected, Loopback0
    B    192.168.100.0/24 [200/0] via 172.16.32.1, 01:45:58
    B*   0.0.0.0/0 [200/0] via 172.16.32.1, 01:45:58
```

SanJose2# show ip route

```
BORIVALI#show ip route
Codes: C - connected, S - static, R - RIP, M - mobile, B -
        D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF
        N1 - OSPF NSSA external type 1, N2 - OSPF NSSA exte
        E1 - OSPF external type 1, E2 - OSPF external type
        i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1,
        ia - IS-IS inter area, * - candidate default, U - p
        o - ODR, P - periodic downloaded static route

Gateway of last resort is 192.168.1.1 to network 0.0.0.0

    172.16.0.0/16 is variably subnetted, 4 subnets, 2 mas
C       172.16.32.0/24 is directly connected, Loopback0
S       172.16.0.0/16 is directly connected, Null0
C       172.16.1.0/24 is directly connected, Serial1/1
D       172.16.64.0/24 [90/2297856] via 172.16.1.1, 02:15:
    192.168.1.0/30 is subnetted, 1 subnets
C       192.168.1.0 is directly connected, Serial1/0
B       192.168.100.0/24 [20/0] via 192.168.1.1, 01:46:10
B*     0.0.0.0/0 [20/0] via 192.168.1.1, 01:46:10
```

- g. Verify the new path using the traceroute command to 10.0.0.1 from SanJose1. Notice the default route is now through SanJose2.

SanJose1# trace 10.0.0.1

```
ANDHERI#trace 10.0.0.1


Type escape sequence to abort.
Tracing the route to 10.0.0.1

 0 172.16.1.2 64 msec 32 msec 32 msec
 1 192.168.1.1 [AS 200] 56 msec 28 msec 64 msec
 2 192.168.1.1 [AS 200] !H !H !H
ANDHERI#
```

Practical No - 4

Aim: Secure the Management Plane

Topology:



Objectives :

- Secure management access.
- Configure enhanced username password security.
- Enable AAA RADIUS authentication.
- Enable secure remote management.

Step 1: Configure loopbacks and assign addresses.

Cable the network as shown in the topology diagram. Erase the startup configuration and reload each router to clear previous configurations. Using the addressing scheme in the diagram, apply the IP addresses to the interfaces on the R1, R2, and R3 routers. You can copy and paste the following configurations into your routers to begin.

Router R1

```
interface Loopback 0
ip address 192.168.1.1 255.255.255.0
```

```
exit
interface Serial0/0/0
  ip address 10.1.1.1 255.255.255.252
  no shutdown
exit
end
```

Router R2

```
interface Serial0/0/0
  ip address 10.1.1.2 255.255.255.252
  no shutdown
exit
interface Serial0/0/1
  ip address 10.2.2.1 255.255.255.252
  no shutdown
exit
end
```

Router R3

```
interface Loopback0
  ip address 192.168.3.1 255.255.255.0
exit
interface Serial0/0/1
  ip address 10.2.2.2 255.255.255.252
  no shutdown
exit
end
```

Step 2: Configure static routes.

```
R1(config)# ip route 0.0.0.0 0.0.0.0 10.1.1.2
```

```
R3(config)# ip route 0.0.0.0 0.0.0.0 10.2.2.1
```

```
R2(config)# ip route 192.168.1.0 255.255.255.0 10.1.1.1
```

```
R2(config)# ip route 192.168.3.0 255.255.255.0 10.2.2.2
```



```
foreach address {  
192.168.1.1  
10.1.1.1  
10.1.1.2  
10.2.2.1  
10.2.2.2  
192.168.3.1  
}{ ping $address }
```

R1# tclsh

```
R1(tcl)#foreach address {  
+>(tcl)#192.168.1.1  
+>(tcl)#10.1.1.1  
+>(tcl)#10.1.1.2  
+>(tcl)#10.2.2.1  
+>(tcl)#10.2.2.2  
+>(tcl)#192.168.3.1  
+>(tcl)# { ping $address }
```

Type escape sequence to abort.

Sending 5, 100-byte ICMP Echos to 192.168.1.1, timeout is 2 seconds:

!!!!

Success rate is 100 percent (5/5), round-trip min/avg/max = 1/1/1 ms Type
escape sequence to abort.

Sending 5, 100-byte ICMP Echos to 10.1.1.1, timeout is 2 seconds:

!!!!

Success rate is 100 percent (5/5), round-trip min/avg/max = 1/2/4 ms Type
escape sequence to abort.

Sending 5, 100-byte ICMP Echos to 10.1.1.2, timeout is 2 seconds:

!!!!

Success rate is 100 percent (5/5), round-trip min/avg/max = 1/1/4 ms Type
escape sequence to abort.

Sending 5, 100-byte ICMP Echos to 10.2.2.1, timeout is 2 seconds:

!!!!

Step 3: Secure management access.

1. On R1, use the **security passwords** command to set a minimum password length of 10 characters.

```
R1(config)# security passwords min-length 10
```

2. Configure the enable secret encrypted password on both routers.

```
R1(config)# enable secret class12345
```

3. Configure a console password and enable login for routers. For additional security, the **exec-timeout** command causes the line to log out after 5 minutes of inactivity. The **logging synchronous** command prevents console messages from interrupting command entry.

```
R1(config)# line console 0
```

```
R1(config-line)# password ciscoconpass
```

```
R1(config-line)# exec-timeout 5 0
```

```
R1(config-line)# login
```

```
R1(config-line)# logging synchronous
```

```
R1(config-line)# exit
```

Configure the password on the vty lines for router R1.

```
R1(config)# line vty 0 4
```

```
R1(config-line)# password ciscovtypass
```

```
R1(config-line)# exec-timeout 5 0
```

```
R1(config-line)# login
```

```
R1(config-line)# exit
```

4. The aux port is a legacy port used to manage a router remotely using a modem and is hardly ever used. Therefore, disable the aux port.

```
R1(config)# line aux 0
```

```
R1(config-line)# no exec
```

```
R1(config-line)# end
```

5. Enter privileged EXEC mode and issue the **show run** command. Can you read the enable secret password? Why or why not?

```
# service password-encryption
```

R1(config)

R1(config)#

6. Configure a warning to unauthorized users with a message-of-the-day (MOTD) banner using the **banner motd** command. When a user connects to one of the routers, the MOTD banner appears before the login prompt. In this example, the dollar sign (\$) is used to start and end the message.

R1(config)# **banner motd \$Unauthorized access strictly prohibited!\$**

R1(config)# **exit**

Step 4: Configure enhanced username password security.

1. To create local database entry encrypted to level 4 (SHA256), use the **username name secret password** global configuration command. In global configuration mode, enter the following command:

R1(config)# **username JR-ADMIN secret class12345**

R1(config)# **username ADMIN secret class54321**

2. Set the console line to use the locally defined login accounts.

R1(config)# **line console 0**

R1(config-line)# **login local**

R1(config-line)# **exit**

3. Set the vty lines to use the locally defined login accounts.

R1(config)# **line vty 0 4**

R1(config-line)# **login local**

R1(config-line)# **end**

4. Repeat the steps 4a to 4c on R3.
 - w. To verify the configuration, telnet to R3 from R1 and login using the ADMIN local database account.

R1# **telnet 10.2.2.2**

Trying 10.2.2.2 ... Open

Unauthorized access strictly prohibited!

User Access Verification

Username: **ADMIN**

Password:

Step 5: Enabling AAA RADIUS Authentication with Local User for Backup.

Configure the specifics for the first RADIUS server located at 192.168.1.101. Use **RADIUS-1-pa55w0rd** as the server password.

```
R1(config)# radius server RADIUS-1
R1(config-radius-server)# address ipv4 192.168.1.101
R1(config-radius-server)# key RADIUS-1-pa55w0rd
R1(config-radius-server)# exit
```

1. Configure the specifics for the second RADIUS server located at 192.168.1.102. Use **RADIUS-2-pa55w0rd** as the server password.

```
R1(config)# radius server RADIUS-2
R1(config-radius-server)# address ipv4 192.168.1.102
R1(config-radius-server)# key RADIUS-2-pa55w0rd
R1(config-radius-server)# exit
```

2. Assign both RADIUS servers to a server group.

```
R1(config)# aaa group server radius RADIUS-GROUP
R1(config-sg-radius)# server name RADIUS-1
R1(config-sg-radius)# server name RADIUS-2
R1(config-sg-radius)# exit
```

3. Enable the default AAA authentication login to attempt to validate against the server group. If they are not available, then authentication should be validated against the local database..

```
R1(config)# aaa authentication login default group RADIUS-GROUP local
```

4. Enable the default AAA authentication Telnet login to attempt to validate against the server group. If they are not available, then authentication should be validated against a case sensitive local database.

```
R1(config)# aaa authentication login TELNET-LOGIN group RADIUS-GROUP local-case
```

Alter the VTY lines to use the TELNET-LOGIN AAA authentication method.

```
R1(config)# line vty 0 4
R1(config-line)# login authentication TELNET-LOGIN
```

R1(config-line)# **exit**

Repeat the steps 5a to 5g on R3.

5. To verify the configuration, telnet to R3 from R1 and login using the ADMIN local database account.

R1# **telnet 10.2.2.2**

Trying 10.2.2.2 ... Open

Unauthorized access strictly prohibited!

User Access Verification

Username: **admin**

Password:

Authentication failed

Username: **ADMIN**

Password:

Step 6: Enabling secure remote management using SSH.

1. SSH requires that a device name and a domain name be configured. Since the router already has a name assigned, configure the domain name.

R1(config)# **ip domain-name ccnasecurity.com**

2. The router uses the RSA key pair for authentication and encryption of transmitted SSH data. Although optional it may be wise to erase any existing key pairs on the router.

R1(config)# **crypto key zeroize rsa**

3. Generate the RSA encryption key pair for the router. Configure the RSA keys with **1024** for the number of modulus bits. The default is 512, and the range is from 360 to 2048.

R1(config)# **crypto key generate rsa general-keys modulus 1024**

The name for the keys will be: R1.ccnasecurity.com

% The key modulus size is 1024 bits

% Generating 1024 bit RSA keys, keys will be non-exportable...[OK]

R1(config)#

Jan 10 13:44:44.711: %SSH-5-ENABLED: SSH 1.99 has been enabled

4. Cisco routers support two versions of SSH:
 - **SSH version 1 (SSHv1)**: Original version but has known vulnerabilities.

- **SSH version 2 (SSHv2):** Provides better security using the Diffie-Hellman key exchange and the strong integrity-checking message authentication code (MAC).

Configure SSH version 2 on R1.

```
R1(config)# ip ssh version 2
```

```
R1(config)#
```

5. Configure the vty lines to use only SSH connections.

```
R1(config)# line vty 0 4
```

```
R1(config-line)# transport input ssh
```

```
R1(config-line)# end
```

6. Verify the SSH configuration using the **show ip ssh** command.

```
R1# show ip ssh
```

```
SSH Enabled - version 2.0
```

```
Authentication timeout: 120 secs; Authentication retries: 3
```

```
Minimum expected Diffie Hellman key size : 1024 bits IOS
```

```
Keys in SECSH format(ssh-rsa, base64 encoded):
```

```
ssh-rsa
```

```
AAAAB3NzaC1yc2EAAAADAQABAAQGC3Lehh7ReYlgyDzls6wq+mFzxqzoaZFr9XGx+Q/ yio
dFYw00hQo80tZy1W1Ff3Pz6q7Qi0y00urwddHZ0kBZceZK9EzJ6wZ+9a87KKDETCWrGSLi6c8lE/y4K+
Z/oVrMMZk7bpTM1MFdP41YgkTf35utYv+TcqbsYo++KjiYk+xw==
```

7. Repeat the steps 6a to 6f on R3.
8. Although a user can SSH from a host using the SSH option of TeraTerm or PuTTY, a router can also SSH to another SSH enabled device. SSH to R3 from R1.

```
R1# ssh -l ADMIN 10.2.2.2
```

```
Password:
```

```
Unauthorized access strictly prohibited!
```

```
R3>
```

```
R3> en
```

```
Password:
```

```
R3#
```

Device Configurations
Router R1

```
service password-encryption
hostname R1
security passwords min-length 10
enable secret 5 $1$t6eK$FZ.JdmMLj8QSgNkpChyZz.
aaa new-model
aaa group server radius RADIUS-GROUP
server name RADIUS-1
server name RADIUS-2
aaa authentication login default group RADIUS-GROUP local
aaa authentication login TELNET-LOGIN group RADIUS-GROUP local-case ip
domain name ccnasecurity.com
username JR-ADMIN secret 5 $1$0u0q$lwimCZlAuQtV4C1ezXL1S0 username
ADMIN secret 5 $1$NSVD$/YjzB7Auyes1sAt4qMfpd.
ip ssh version 2
interface Loopback0
description R1 LAN
ip address 192.168.1.1 255.255.255.0
interface Serial0/0/0
description R1 --> R2
ip address 10.1.1.1 255.255.255.252
no fair-queue
ip route 0.0.0.0 0.0.0.0 10.1.1.2
radius server RADIUS-1
address ipv4 192.168.1.101 auth-port 1645 acct-port 1646
key 7 107C283D2C2221465D493A2A717D24653017
radius server RADIUS-2
address ipv4 192.168.1.102 auth-port 1645 acct-port 1646
key 7 03367A2F2F3A12011C44090442471C5C162E
banner motd ^CUnauthorized access strictly prohibited!^C

line con 0
exec-timeout 5 0
```

```
password 7 070C285F4D061A0A19020A1F17
```

```
logging synchronous
```

```
line aux 0
```

```
no exec
```

```
password 7 060506324F411F0D1C0713181F
```

```
login authentication TELNET-LOGIN
```

```
transport input ssh
```

```
end
```

Router R2

```
hostname R2
```

```
enable secret 5 $1$DJS7$xvJDW87zLs8pSJDFULCPB1
```

```
interface Serial0/0/0
```

```
ip address 10.1.1.2 255.255.255.252
```

```
no fair-queue
```

```
interface Serial0/0/1
```

```
ip address 10.2.2.1 255.255.255.252
```

```
clock rate 128000
```

```
ip route 192.168.1.0 255.255.255.0 10.1.1.1
```

```
ip route 192.168.3.0 255.255.255.0 10.2.2.2
```

```
line con 0
```

```
exec-timeout 0 0
```

```
logging synchronous
```

```
line vty 0 4
```

```
password cisco
```

```
login
```

```
end
```

Router R3

```
service password-encryption
```



```
hostname R3
security passwords min-length 10
enable secret 5 $1$50Y4$4J6VFlvGNKjwQ8XtajgUk1
aaa new-model
aaa group server radius RADIUS-GROUP
server name RADIUS-1
server name RADIUS-2
aaa authentication login default group RADIUS-GROUP local
aaa authentication login TELNET-LOGIN group RADIUS-GROUP local-case ip
domain name ccnasecurity.com
```

```
username JR-ADMIN secret 5 $1$b4m1$RVmjL9S3gxKh1xr8qzNqr/
username ADMIN secret 5 $1$zGV7$pVgSEbinvXQ7f7uyxeKBj
ip ssh version 2
```

```
interface Loopback0
description R3 LAN
ip address 192.168.3.1 255.255.255.0
```

```
interface Serial0/0/1
description R3 --> R2
ip address 10.2.2.2 255.255.255.252
```

```
ip route 0.0.0.0 0.0.0.0 10.2.2.1
```

```
radius server RADIUS-1
address ipv4 192.168.1.101 auth-port 1645 acct-port 1646
key 7 01212720723E354270015E084C5000421908
```

```
radius server RADIUS-2
address ipv4 192.168.1.102 auth-port 1645 acct-port 1646
```

```
key 7 003632222D6E384B5D6C5C4F5C4C1247000F
```

```
banner motd ^CUnauthorized access strictly prohibited!^C
```

```
line con 0
```

```
exec-timeout 5 0
```

```
password 7 104D000A0618110402142B3837
```

```
logging synchronous
```

```
line aux 0
```

```
no exec
```

```
line vty 0 4
```

```
exec-timeout 5 0
```

```
password 7 070C285F4D060F110E020A1F17
```

```
login authentication TELNET-LOGIN
```

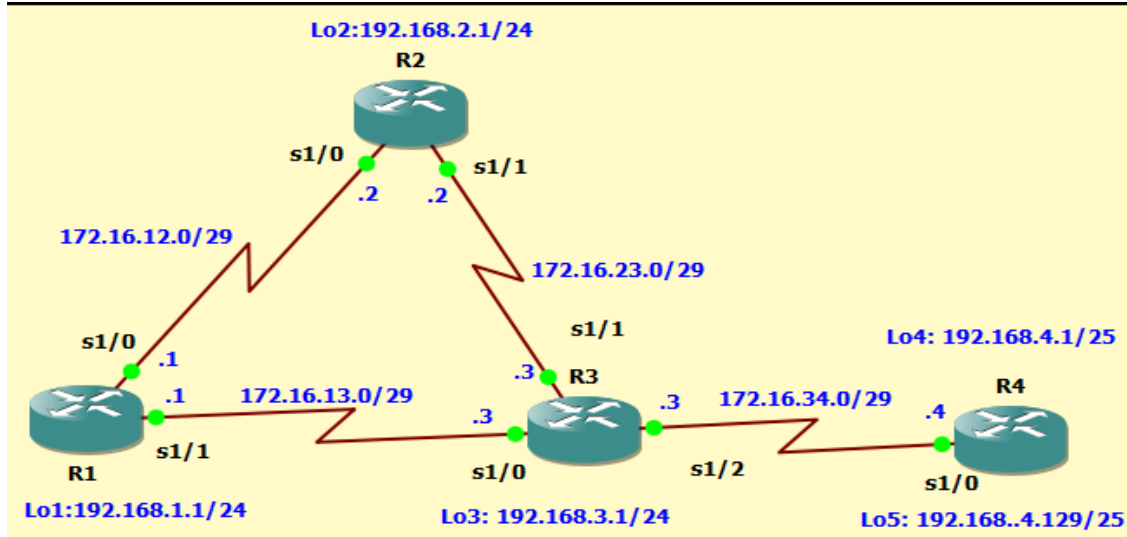
```
transport input ssh
```

```
end
```

Practical No - 5

Aim : Configure and Verify Path Control Using PBR

Topology :



Objectives

- Configure and verify policy-based routing.
- Select the required tools and commands to configure policy-based routing operations.
- Verify the configuration and operation by using the proper show and debug commands.

Step 1: Configure loopbacks and assign addresses.

- Cable the network as shown in the topology diagram. Erase the startup configuration, and reload each router to clear previous configurations.
- Using the addressing scheme in the diagram, create the loopback interfaces and apply IP addresses to these and the serial interfaces on R1, R2, R3, and R4. On the serial interfaces connecting R1 to R3 and R3 to R4, specify the bandwidth as 64 Kb/s and set a clock rate on the DCE using the **clock rate 64000** command. On the serial interfaces connecting R1 to R2 and R2 to R3, specify the bandwidth as 128 Kb/s and set a clock rate on the DCE using the **clock rate 128000** command.

You can copy and paste the following configurations into your routers to begin.

Note: Depending on the router model, interfaces might be numbered differently than those listed. You might need to alter them accordingly.

Router R1

```
interface Lo1
ip address 192.168.1.1 255.255.255.0
interface Serial0/0/0
ip address 172.16.12.1 255.255.255.248
no shutdown
```

```
interface Serial0/0/1
ip address 172.16.13.1 255.255.255.248
no shutdown
End
```

```
R1(config)#int Lo1
```

```
R1(config-if)#ip address 192.168.1.1 255.255.255.0
R1(config-if)#int s1/0
R1(config-if)#ip address 172.16.12.1 255.255.255.248
R1(config-if)#no shutdown
```

```
R1(config-if)#int s1/1
R1(config-if)#ip address 172.16.13.1 255.255.255.248
R1(config-if)#no shutdown
*May 19 23:06:21.987: %LINEPROTO-5-UPDOWN: Line protocol is down
R1(config-if)#no shutdown
```

Router R2

```
interface Lo2
ip address 192.168.2.1 255.255.255.0
interface Serial0/0/0
ip address 172.16.12.2 255.255.255.248
no shutdown
```

```
interface Serial0/0/1
ip address 172.16.23.2 255.255.255.248
no shutdown
End
```

```
R2(config)#int Lo2
R2(config-if)#
*May 19 23:06:13.083: %LINEPROTO-5-UPDOWN: Line protocol is down
R2(config-if)#ip address 192.168.2.1 255.255.255.0
R2(config-if)#int s1/0
R2(config-if)#ip address 172.16.12.2 255.255.255.248
R2(config-if)#no shutdown
```

```
R2(config)#int s1/1
R2(config-if)#ip address 172.16.23.2 255.255.255.248
R2(config-if)#no shutdown
```

Router R3

```
interface Lo3
ip address 192.168.3.1 255.255.255.0
```

```
interface Serial0/0/0
  ip address 172.16.13.3 255.255.255.248
  no shutdown
interface Serial0/0/1
  ip address 172.16.23.3 255.255.255.248
  no shutdown
interface Serial0/1/0
  ip address 172.16.34.3 255.255.255.248
  no shutdown
End
```

```
R3(config)#int Lo3
R3(config-if)#
*May 19 23:07:08.351: %LINEPROTO-5-UPDOWN: Line protocol is down
R3(config-if)#ip address 192.168.3.1 255.255.255.0
R3(config-if)#int s1/0
R3(config-if)#ip address 172.16.13.3 255.255.255.248
R3(config-if)#no shutdown
```

```
R3(config-if)#int s1/1
R3(config-if)#ip address 172.16.23.3 255.255.255.248
R3(config-if)#no shutdown
```

```
R3(config-if)#int s1/2
R3(config-if)#ip address 172.16.34.3 255.255.255.248
R3(config-if)#no shutdown
R3(config-if)#exit
```

Router R4

```
interface Lo4
  ip address 192.168.4.1 255.255.255.128
interface Lo5
  ip address 192.168.4.129 255.255.255.128
interface Serial0/0/0
  ip address 172.16.34.4 255.255.255.248
  no shutdown
End
```

```

R4(config)#int lo4
R4(config-if)#
*May 19 23:08:16.239: %LINEPROTO-5-UPDOWN: Line protocol is down
R4(config-if)#ip address 192.168.4.1 255.255.255.128
R4(config-if)#interface Lo5
R4(config-if)#
*May 19 23:08:32.527: %LINEPROTO-5-UPDOWN: Line protocol is down
R4(config-if)#ip address 192.168.4.129 255.255.255.128
R4(config-if)#int s1/0
R4(config-if)#ip address 172.16.34.4 255.255.255.248
R4(config-if)#no shutdown

```

- z. Verify the configuration with the **show ip interface brief**, **show protocols**, and **show interfaces description** commands. The output from router R3 is shown here as an example.

R3# show ip interface brief

```

R3#show ip interface brief
Interface                IP-Address      OK? Method Status                Protocol
FastEthernet0/0          unassigned      YES unset  administratively down  down
Serial1/0                 172.16.13.3     YES manual  up                    up
Serial1/1                 172.16.23.3     YES manual  up                    up
Serial1/2                 172.16.34.3     YES manual  up                    up
Serial1/3                 unassigned      YES unset  administratively down  down
Loopback3                 192.168.3.1     YES manual  up                    up

```

R3# show protocols

```

R3#show protocols
Global values:
  Internet Protocol routing is enabled
FastEthernet0/0 is administratively down, line protocol is down
Serial1/0 is up, line protocol is up
  Internet address is 172.16.13.3/29
Serial1/1 is up, line protocol is up
  Internet address is 172.16.23.3/29
Serial1/2 is up, line protocol is up
  Internet address is 172.16.34.3/29
Serial1/3 is administratively down, line protocol is down
Loopback3 is up, line protocol is up
  Internet address is 192.168.3.1/24

```

R3# show interfaces description

```

R3#show interfaces description
Interface                Status          Protocol Description
Fa0/0                    admin down      down
Se1/0                    up              up
Se1/1                    up              up
Se1/2                    up              up
Se1/3                    admin down      down

```

Step 3: Configure basic EIGRP.

- aa. Implement EIGRP AS 1 over the serial and loopback interfaces as you have configured it for the other EIGRP labs.
- bb. Advertise networks 172.16.12.0/29, 172.16.13.0/29, 172.16.23.0/29, 172.16.34.0/29, 192.168.1.0/24, 192.168.2.0/24, 192.168.3.0/24, and 192.168.4.0/24 from their respective routers.

You can copy and paste the following configurations into your routers.

Router R1

```
router eigrp 1
network 192.168.1.0
network 172.16.12.0 0.0.0.7
network 172.16.13.0 0.0.0.7
no auto-summary
```

```
R1(config)#router eigrp 1
R1(config-router)#network 192.168.1.0
R1(config-router)#network 172.16.12.0 0.0.0.7
R1(config-router)#network 172.16.13.0 0.0.0.7
R1(config-router)#no auto-summary
```

Router R2

```
router eigrp 1
network 192.168.2.0
network 172.16.12.0 0.0.0.7
network 172.16.23.0 0.0.0.7
no auto-summary
```

```
R2(config)#
R2(config)#router eigrp 1
R2(config-router)#network 192.168.2.0
R2(config-router)#network 172.16.12.0 0.0.0.7
R2(config-router)#network 172.16.23.0 0.0.0.7
R2(config-router)#no auto-summary
```

Router R3

```
router eigrp 1
network 192.168.3.0
network 172.16.13.0 0.0.0.7
network 172.16.23.0 0.0.0.7
network 172.16.34.0 0.0.0.7
no auto-summary
```

```
R3(config)#router eigrp 1
R3(config-router)#network 192.168.3.0
R3(config-router)#network 172.16.13.0 0.0.0.7
R3(config-router)#network 172.16.23.0 0.0.0.7
R3(config-router)#network 172.16.34.0 0.0.0.7
R3(config-router)#no auto-summary
```

Router R4

```
router eigrp 1
```

```
network 192.168.4.0
```

```
network 172.16.34.0 0.0.0.7
```

```
no auto-summary
```

```
R4(config)#router eigrp 1
R4(config-router)#network 192.168.4.0
R4(config-router)#network 172.16.34.0 0.0.0.7
R4(config-router)#no auto-summary
```

You should see EIGRP neighbor relationship messages being generated.

Step 4: Verify EIGRP connectivity.

- cc. Verify the configuration by using the **show ip eigrp neighbors** command to check which routers have EIGRP adjacencies.

R1# show ip eigrp neighbors

```
R1#show ip eigrp neighbors
IP-EIGRP neighbors for process 1
H   Address          Interface      Hold Uptime    SRTT  RT0  Q  Seq
   (sec)              (ms)          Cnt  Num
1   172.16.13.3        Se1/1         11 00:00:31    26   200  0  18
0   172.16.12.2        Se1/0         12 00:00:44    37   222  0  13
```

R2# show ip eigrp neighbors

```
R2#show ip eigrp neighbors
*May 19 23:13:42.783: %SYS-5-CONFIG_I: Configured from console by console
R2#show ip eigrp neighbors
IP-EIGRP neighbors for process 1
H   Address          Interface      Hold Uptime    SRTT  RT0  Q  Seq
   (sec)              (ms)          Cnt  Num
1   172.16.23.3        Se1/1         11 00:00:50    41   246  0  20
0   172.16.12.1        Se1/0         11 00:01:04    30   200  0  18
```

R3# show ip eigrp neighbors

```
R3#show ip eigrp neighbors
IP-EIGRP neighbors for process 1
H   Address          Interface      Hold Uptime    SRTT  RT0  Q  Seq
   (sec)              (ms)          Cnt  Num
2   172.16.34.4        Se1/2         12 00:00:44    48   288  0  6
1   172.16.23.2        Se1/1         11 00:00:58    26   200  0  19
0   172.16.13.1        Se1/0         12 00:00:58    281  1686  0  20
```

R4# show ip eigrp neighbors

```
R4#show ip eigrp neighbors
IP-EIGRP neighbors for process 1
H   Address          Interface      Hold Uptime    SRTT  RT0  Q  Seq
   (sec)              (ms)          Cnt  Num
0   172.16.34.3        Se1/0         10 00:00:55    23   200  0  26
```

- dd. Run the following Tcl script on all routers to verify full connectivity.

R1# tclsh

```
R1#tclsh
R1(tcl)#foreach address {
+>(tcl)#172.16.12.1
+>(tcl)#172.16.12.2
+>(tcl)#172.16.13.1
+>(tcl)#172.16.13.3
+>(tcl)#172.16.23.2
+>(tcl)#172.16.23.3
+>(tcl)#172.16.34.3
+>(tcl)#172.16.34.4
+>(tcl)#192.168.1.1
+>(tcl)#192.168.2.1
+>(tcl)#192.168.3.1
+>(tcl)#192.168.4.1
+>(tcl)#192.168.4.129
+>(tcl)#} { ping $address }
```

```
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 172.16.12.1, timeout is 2 seconds:
!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 44/61/76 ms
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 172.16.12.2, timeout is 2 seconds:
!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 16/27/40 ms
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 172.16.13.1, timeout is 2 seconds:
!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 40/58/80 ms
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 172.16.13.3, timeout is 2 seconds:
!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 20/31/44 ms
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 172.16.23.2, timeout is 2 seconds:
!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 20/28/32 ms
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 172.16.23.3, timeout is 2 seconds:
!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 24/40/48 ms
Type escape sequence to abort.
```

Step 5: Verify the current path.

Before you configure PBR, verify the routing table on R1.

- ee. On R1, use the **show ip route** command. Notice the next-hop IP address for all networks discovered by EIGRP.

R1# show ip route

```

R1#show ip route
Codes: C - connected, S - static, R - RIP, M - mobile, B - BGP
       D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter area
       N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type 2
       E1 - OSPF external type 1, E2 - OSPF external type 2
       i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1, L2 - IS-IS le
       ia - IS-IS inter area, * - candidate default, U - per-user stati
       o - ODR, P - periodic downloaded static route

Gateway of last resort is not set

    172.16.0.0/29 is subnetted, 4 subnets
O       172.16.34.0 [90/2681856] via 172.16.13.3, 00:01:50, Serial1/1
O       172.16.23.0 [90/2681856] via 172.16.13.3, 00:01:50, Serial1/1
        [90/2681856] via 172.16.12.2, 00:01:50, Serial1/0
C       172.16.12.0 is directly connected, Serial1/0
C       172.16.13.0 is directly connected, Serial1/1
O     192.168.4.0/24 [90/2809856] via 172.16.13.3, 00:01:38, Serial1/1
C     192.168.1.0/24 is directly connected, Loopback1
O     192.168.2.0/24 [90/2297856] via 172.16.12.2, 00:01:50, Serial1/0
O     192.168.3.0/24 [90/2297856] via 172.16.13.3, 00:01:50, Serial1/1

```

R4# traceroute 192.168.1.1 source 192.168.4.1

```

R4#traceroute 192.168.1.1 source 192.168.4.1

Type escape sequence to abort.
Tracing the route to 192.168.1.1

  1 172.16.34.3 36 msec 32 msec 32 msec
  2 172.16.13.1 28 msec 56 msec 84 msec

```

R4# traceroute 192.168.1.1 source 192.168.4.129

```

R4#traceroute 192.168.1.1 source 192.168.4.129

Type escape sequence to abort.
Tracing the route to 192.168.1.1

  1 172.16.34.3 44 msec 28 msec 28 msec
  2 172.16.13.1 64 msec 28 msec 64 msec

```

On R3, use the **show ip route** command and note that the preferred route from R3 to R1 LAN 192.168.1.0/24 is via R2 using the R3 exit interface S0/0/1.

R3# show ip route

```

R3#show ip route
Codes: C - connected, S - static, R - RIP, M - mobile, B - BGP
       D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter area
       N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type 2
       E1 - OSPF external type 1, E2 - OSPF external type 2
       i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1, L2 - IS-IS leve
       ia - IS-IS inter area, * - candidate default, U - per-user static
       o - ODR, P - periodic downloaded static route

Gateway of last resort is not set

```

```

172.16.0.0/29 is subnetted, 4 subnets
C    172.16.34.0 is directly connected, Serial1/2
C    172.16.23.0 is directly connected, Serial1/1
D    172.16.12.0 [90/2681856] via 172.16.23.2, 00:02:29, Serial1/1
      [90/2681856] via 172.16.13.1, 00:02:29, Serial1/0
C    172.16.13.0 is directly connected, Serial1/0
D    192.168.4.0/24 [90/2297856] via 172.16.34.4, 00:02:17, Serial1/2
D    192.168.1.0/24 [90/2297856] via 172.16.13.1, 00:02:29, Serial1/0
D    192.168.2.0/24 [90/2297856] via 172.16.23.2, 00:02:29, Serial1/1
C    192.168.3.0/24 is directly connected, Loopback3

```

R3#

ff. On R3, use the **show interfaces serial 0/0/0** and **show interfaces s0/0/1** commands.

R3# **show interfaces serial0/0/0**

```

R3#show int s1/0
Serial1/0 is up, line protocol is up
  Hardware is M4T
  Internet address is 172.16.13.3/29
  MTU 1500 bytes, BW 1544 Kbit/sec, DLY 20000 usec,

```

```

Routing Descriptor Blocks:
  172.16.13.1 (Serial1/0), from 172.16.13.1, Send flag is 0x0
    Composite metric is (2297856/128256), Route is Internal
    Vector metric:
      Minimum bandwidth is 1544 Kbit
      Total delay is 25000 microseconds
      Reliability is 255/255
      Load is 1/255
      Minimum MTU is 1500
      Hop count is 1
  172.16.23.2 (Serial1/1), from 172.16.23.2, Send flag is 0x0

```

gg. Confirm that R3 has a valid route to reach R1 from its serial 0/0/0 interface using the **show ip eigrp topology 192.168.1.0** command.

R3# **show ip eigrp topology 192.168.1.0**

```

R3#show ip eigrp topology 192.168.1.0
IP-EIGRP (AS 1): Topology entry for 192.168.1.0/24
  State is Passive, Query origin flag is 1, 1 Successor(s), FD is
  Routing Descriptor Blocks:
    172.16.13.1 (Serial1/0), from 172.16.13.1, Send flag is 0x0
      Composite metric is (2297856/128256), Route is Internal
      Vector metric:
        Minimum bandwidth is 1544 Kbit
        Total delay is 25000 microseconds
        Reliability is 255/255
        Load is 1/255
        Minimum MTU is 1500
        Hop count is 1
    172.16.23.2 (Serial1/1), from 172.16.23.2, Send flag is 0x0

```

Step 6: Configure PBR to provide path control.

The steps required to implement path control include the following:

- Choose the path control tool to use. Path control tools manipulate or bypass the IP routing table. For PBR, **route-map** commands are used.
- Implement the traffic-matching configuration, specifying which traffic will be manipulated. The **match** commands are used within route maps.
- Define the action for the matched traffic using **set** commands within route maps.
- Apply the route map to incoming traffic.

As a test, you will configure the following policy on router R3:

- All traffic sourced from R4 LAN A must take the R3 --> R2 --> R1 path.
 - All traffic sourced from R4 LAN B must take the R3 --> R1 path.
- hh. On router R3, create a standard access list called **PBR-ACL** to identify the R4 LAN B network.

```
R3(config)# ip access-list standard PBR-ACL
```

```
R3(config-std-nacl)# remark ACL matches R4 LAN B traffic
```

```
R3(config-std-nacl)# permit 192.168.4.128 0.0.0.127
```

```
R3(config-std-nacl)# exit
```

```
R3(config)#ip access-list standard PBR-ACL
R3(config-std-nacl)#remark ACL matches R4 LAN B traffic
R3(config-std-nacl)#permit 192.168.4.128 0.0.0.127
R3(config-std-nacl)#exit
```

```
R3(config)#
```

- ii. Create a route map called **R3-to-R1** that matches PBR-ACL and sets the next-hop interface to the R1 serial 0/0/1 interface.

```
R3(config)# route-map R3-to-R1 permit
```

```
R3(config-route-map)# description RM to forward LAN B traffic to R1
```

```
R3(config-route-map)# match ip address PBR-ACL
```

```
R3(config-route-map)# set ip next-hop 172.16.13.1
```

```
R3(config-route-map)# exit
```

```
R3(config)#route-map R3-to-R1 permit
R3(config-route-map)#match ip address PBR-ACL
R3(config-route-map)#set ip next-hop 172.16.13.1
R3(config-route-map)#exit
```

- jj. Apply the R3-to-R1 route map to the serial interface on R3 that receives the traffic from R4. Use the **ip policy route-map** command on interface S0/1/0.

```
R3(config)# interface s0/1/0
```

```
R3(config-if)# ip policy route-map R3-to-R1
```

```
R3(config-if)# end
```

```
R3(config)#int s1/2
R3(config-if)#ip policy route-map R3-to-R1
R3(config-if)#end
```

kk. On R3, display the policy and matches using the **show route-map** command.

R3# **show route-map**

```
R3#show route-map
route-map R3-to-R1, permit, sequence 10
  Match clauses:
    ip address (access-lists): PBR-ACL
  Set clauses:
    ip next-hop 172.16.13.1
  Policy routing matches: 0 packets, 0 bytes
R3#
```

Step 7: Test the policy.

ll. On R3, create a standard ACL which identifies all of the R4 LANs.

R3# **conf t**

Enter configuration commands, one per line. End with CNTL/Z. R3(config)#

access-list 1 permit 192.168.4.0 0.0.0.255

R3(config)# **exit**

```
R3#conf t
Enter configuration commands, one per line. End with CNTL/Z.
R3(config)#access-list 1 permit 192.168.4.0 0.0.0.255
R3(config)#exit
```

mm. Enable PBR debugging only for traffic that matches the R4 LANs.

R3# **debug ip policy ?**

```
R3#debug ip policy ?
<1-199> Access list
dynamic dynamic PBR
<cr>
```

R3# **debug ip policy 1**

```
R3#debug ip policy 1
Policy routing debugging is on for access list 1
```

nn. Test the policy from R4 with the **traceroute** command, using R4 LAN A as the source network.

R4# **traceroute 192.168.1.1 source 192.168.4.1**

```
R4#traceroute 192.168.1.1 source 192.168.4.1

Type escape sequence to abort.
Tracing the route to 192.168.1.1

 0 172.16.34.3 40 msec 12 msec 32 msec
 1 172.16.13.1 60 msec 48 msec 88 msec
```

```

R3#
*May 19 23:17:36.819: IP: s=192.168.4.1 (Serial1/2), d=192.168.1.1,
*May 19 23:17:36.851: IP: s=192.168.4.1 (Serial1/2), d=192.168.1.1,
*May 19 23:17:36.879: IP: s=192.168.4.1 (Serial1/2), d=192.168.1.1,
*May 19 23:17:36.915: IP: s=192.168.4.1 (Serial1/2), d=192.168.1.1,
S
*May 19 23:17:36.971: IP: s=192.168.4.1 (Serial1/2), d=192.168.1.1,
S
*May 19 23:17:37.031: IP: s=192.168.4.1 (Serial1/2), d=192.168.1.1
R3#, len 28, FIB policy rejected(no match) - normal forwarding
R3#

```

oo. Test the policy from R4 with the **traceroute** command, using R4 LAN B as the source network.

R4# **traceroute 192.168.1.1 source 192.168.4.129**

```

R4#traceroute 192.168.1.1 source 192.168.4.129

Type escape sequence to abort.
Tracing the route to 192.168.1.1

 0 172.16.34.3 40 msec 28 msec 32 msec
 1 172.16.13.1 60 msec 64 msec 32 msec

```

```

R3#.168.4.129 (Serial1/2), d=192.168.1.1, len 28, FIB policy match
*May 19 23:17:55.763: IP: s=192.168.4.129 (Serial1/2), d=192.168.1.1,
*May 19 23:17:55.763: IP: s=192.168.4.129 (Serial1/2), d=192.168.1.1,
*May 19 23:17:55.823: IP: s=192.168.4.129 (Serial1/2), d=192.168.1.1,
*May 19 23:17:55.823: IP: s=192.168.4.129 (Serial1/2), d=192.168.1.1,
*May 19 23:17:55.827: IP: s=192.168.4.129 (Serial1/2), d=192.168.1.1,
*May 19 23:17:55.883: IP: s=192.168.4.129 (Serial1/2), d=192.168.1.1,
*May 19 23:17:55.883: IP: s=192.168.4.129 (Serial1/2), d=192.168.1.1,
*May 19 23:17:55.887: IP: s=192.168.4.129 (Serial1/2), d=192.168.1.1,

```

pp. On R3, display the policy and matches using the **show route-map** command.

R3# **show route-map**

```

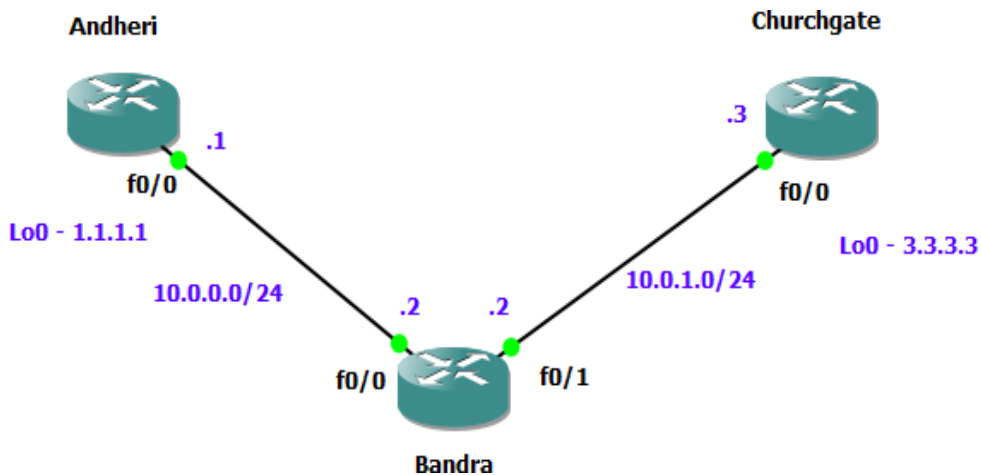
R3#show route-map
route-map R3-to-R1, permit, sequence 10
  Match clauses:
    ip address (access-lists): PBR-ACL
  Set clauses:
    ip next-hop 172.16.13.1
  Policy routing matches: 6 packets, 192 bytes

```

Practical No - 6

Aim: Cisco MPLS Configuration

Topology :



Step 1 - IP addressing of MPLS Core and OSPF

First bring 3 routers into your topology R1, R2, R3 position them as below. We are going to address the routers and configure ospf to ensure loopback to loopback connectivity between R1 and R3

```
Andheri(config)#int lo0
Andheri(config-if)#ip add 1.1.1.1 255.255.255.255
Andheri(config-if)#ip ospf 1 area 0
Andheri(config-if)#
Andheri(config-if)#int f0/0
Andheri(config-if)#ip add 10.0.0.1 255.255.255.0
Andheri(config-if)#no shut
Andheri(config-if)#ip ospf 1 area 0
```



```

Bandra(config)#int lo0
Bandra(config-if)#
Bandra(config-if)#ip add 2.2.2.2 255.255.255.255
Bandra(config-if)#ip ospf 1 area 0
Bandra(config-if)#
Bandra(config-if)#int f0/0
Bandra(config-if)#ip add 10.0.0.2 255.255.255.0
Bandra(config-if)#no shut
Bandra(config-if)#ip ospf 1 area 0
Bandra(config-if)#
Bandra(config-if)#int f0/1
Bandra(config-if)#ip add 10.0.1.2 255.255.255.0
Bandra(config-if)#no shut
Bandra(config-if)#ip ospf 1 area 0

```

```

Churchgate(config)#int lo0
Churchgate(config-if)#ip add 3.3.3.3 255.255.255.255
Churchgate(config-if)#ip ospf 1 area 0
Churchgate(config-if)#
Churchgate(config-if)#int f0/0
Churchgate(config-if)#ip add 10.0.1.3 255.255.255.0
Churchgate(config-if)#no shut
Churchgate(config-if)#ip ospf 1 area 0

```

You should now have full ip connectivity between R1, R2, R3 to verify this we need to see if we can ping between the loopbacks of R1 and R3

```

Andheri#ping 3.3.3.3 source lo0
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 3.3.3.3, timeout is 2 seconds:
Packet sent with a source address of 1.1.1.1
!!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 20/52/64 ms

```

Step 2 - Configure LDP on all the interfaces in the MPLS Core

In order to run MPLS you need to enable it, there are two ways to do this.

- ☐ At each interface enter the mpls ip command
- ☐ Under the ospf process use the mpls ldp autoconfig command

```

Andheri(config)#router ospf 1
Andheri(config-router)#mpls ldp autoconfig

```

```

Bandra(config)#router ospf 1
Bandra(config-router)#mpls ldp autoconfig

```

```

Churchgate(config)#router ospf 1
Churchgate(config-router)#mpls ldp autoconfig

```

You should see log messages coming up showing the LDP neighbors are up.

```

Bandra#
*May 29 17:03:09.559: %SYS-5-CONFIG_I: Configured from console by console
Bandra#
*May 29 17:03:28.631: %LDP-5-NBRCHG: LDP Neighbor 3.3.3.3:0 (2) is UP

```

To verify the mpls interfaces the command is very simple – sh mpls interface

This is done on R2 and you can see that both interfaces are running mpls and using LDP

```

Bandra#sh mpls int
Interface          IP          Tunnel  BGP Static Operational
FastEthernet0/0    Yes (ldp)   No      No  No   Yes
FastEthernet0/1    Yes (ldp)   No      No  No   Yes
Bandra#

```

You can also verify the LDP neighbors with the sh mpls ldp neighbors command.

```

Bandra#sh mpls ldp neigh
  Peer LDP Ident: 1.1.1.1:0; Local LDP Ident 2.2.2.2:0
    TCP connection: 1.1.1.1.646 - 2.2.2.2.25712
    State: Oper; Msgs sent/rcvd: 9/9; Downstream
    Up time: 00:01:23
    LDP discovery sources:
      FastEthernet0/0, Src IP addr: 10.0.0.1
    Addresses bound to peer LDP Ident:
      10.0.0.1      1.1.1.1
  Peer LDP Ident: 3.3.3.3:0; Local LDP Ident 2.2.2.2:0
    TCP connection: 3.3.3.3.50470 - 2.2.2.2.646
    State: Oper; Msgs sent/rcvd: 8/8; Downstream
    Up time: 00:00:54
    LDP discovery sources:
      FastEthernet0/1, Src IP addr: 10.0.1.3
    Addresses bound to peer LDP Ident:
      10.0.1.3      3.3.3.3

```

One more verification to confirm LDP is running ok is to do a trace between R1 and R3 and verify if you get MPLS Labels show up in the trace.

```

Andheri#trace 3.3.3.3

Type escape sequence to abort.
Tracing the route to 3.3.3.3

  1 10.0.0.2 [MPLS: Label 17 Exp 0] 20 msec 60 msec 60 msec
  2 10.0.1.3 60 msec 60 msec 60 msec

```

Step 3 - MPLS BGP Configuration between R1 and R3

We need to establish a Multi Protocol BGP session between R1 and R3 this is done by configuring the vpnv4 address family as below

```

Andheri(config)#router bgp 1
Andheri(config-router)#neighbor 3.3.3.3 remote-as 1
Andheri(config-router)#neighbor 3.3.3.3 update-source Loopback0
Andheri(config-router)#no auto-summary
Andheri(config-router)#!
Andheri(config-router)#address-family vpnv4
Andheri(config-router-af)#neighbor 3.3.3.3 activate

```

```

Churchgate(config)#router bgp 1
Churchgate(config-router)#neighbor 1.1.1.1 remote-as 1
Churchgate(config-router)#neighbor 1.1.1.1
*May 29 17:06:19.459: %BGP-5-ADJCHANGE: neighbor 1.1.1.1 Up
Churchgate(config-router)#neighbor 1.1.1.1 update-source loopback 0
Churchgate(config-router)#no auto-summary
Churchgate(config-router)#address-family vpnv4
Churchgate(config-router-af)#neighbor 1.1.1.1 activate

```

To verify the BGP session between R1 and R3 issue the command `sh bgp vpnv4 unicast all summary`

```

Andheri#sh bgp vpnv4 unicast all summary
BGP router identifier 1.1.1.1, local AS number 1
BGP table version is 1, main routing table version 1

Neighbor      V      AS MsgRcvd MsgSent  TblVer  InQ  OutQ Up/Down  State/PfxRcd
3.3.3.3        4        1      5       6       1    0    0 00:00:30 0/0

```

Step 4 - Add two more routers, create VRFs

We will add two more routers into the topology so it now looks like the final topology

```

Borivali(config)#int lo0
Borivali(config-if)#ip ad
*May 29 17:13:47.223: %LINEPROTO-5-UPDOWN: Line protocol on Interface Loopback0
Borivali(config-if)#ip address 4.4.4.4 255.255.255.255
Borivali(config-if)#ip ospf 2 area 2
Borivali(config-if)#int f0/0
Borivali(config-if)#ip address 192.168.1.4 255.255.255.0
                                     ^
% Invalid input detected at '^' marker.

Borivali(config-if)#ip address 192.168.1.4 255.255.255.0
Borivali(config-if)#ip ospf 2 area 2
Borivali(config-if)#no shut

```

```

Andheri(config)#int f0/1
Andheri(config-if)#no shut
Andheri(config-if)#ip address
*May 29 17:14:16.199: %LINK-3-UPDOWN: Interface FastEthernet0/1 is up
*May 29 17:14:17.199: %LINEPROTO-5-UPDOWN: Line protocol on Interface FastEthernet0/1
Andheri(config-if)#ip address 192.168.1.1 255.255.255.0

```

```

Andheri(config-if)#ip vrf RED
Andheri(config-vrf)#rd 4:4
Andheri(config-vrf)#route-target both 4:4

```

```

Andheri(config-vrf)#int f0/1
Andheri(config-if)#ip vrf forwarding RED
% Interface FastEthernet0/1 IP address 192.168.1.1 removed due to enabling VRF RED

```

```

Andheri#sh run int f0/1
Building configuration...

Current configuration : 119 bytes
!
interface FastEthernet0/1
 ip vrf forwarding RED
 ip address 192.168.1.1 255.255.255.0
 duplex auto
 speed auto
end

```

If you issue the command `sh ip route` this shows the routes in the global table and you will notice that you do not see 192.168.1.0/24

```

Andheri#sh ip route
Codes: C - connected, S - static, R - RIP, M - mobile, B - BGP
       D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter area
       N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type 2
       E1 - OSPF external type 1, E2 - OSPF external type 2
       i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1, L2 - IS-IS level-2
       ia - IS-IS inter area, * - candidate default, U - per-user static route
       o - ODR, P - periodic downloaded static route

Gateway of last resort is not set

 1.0.0.0/32 is subnetted, 1 subnets
C       1.1.1.1 is directly connected, Loopback0
 2.0.0.0/32 is subnetted, 1 subnets
O       2.2.2.2 [110/2] via 10.0.0.2, 00:19:39, FastEthernet0/0
 3.0.0.0/32 is subnetted, 1 subnets
O       3.3.3.3 [110/3] via 10.0.0.2, 00:18:35, FastEthernet0/0
10.0.0.0/24 is subnetted, 2 subnets
C       10.0.0.0 is directly connected, FastEthernet0/0
O       10.0.1.0 [110/2] via 10.0.0.2, 00:18:45, FastEthernet0/0

```

```

Andheri#sh ip route vrf RED
Routing Table: RED
Codes: C - connected, S - static, R - RIP, M - mobile, B - BGP
       D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter area
       N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type 2
       E1 - OSPF external type 1, E2 - OSPF external type 2
       i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1, L2 - IS-IS level-2
       ia - IS-IS inter area, * - candidate default, U - per-user static route
       o - ODR, P - periodic downloaded static route

Gateway of last resort is not set

C       192.168.1.0/24 is directly connected, FastEthernet0/1

```

We just need to enable OSPF on this interface and get the loopback address for R4 in the VRF RED routing table before proceeding.

```

Andheri(config)#int f0/1
Andheri(config-if)#ip ospf 2 area 2

```

If we now check the routes in the VRF RED routing table you should see 4.4.4.4 in there as well.

```
Andheri#sh ip route vrf RED

Routing Table: RED
Codes: C - connected, S - static, R - RIP, M - mobile, B - BGP
       D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter ar
       N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type
       E1 - OSPF external type 1, E2 - OSPF external type 2
       i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1, L2 - IS-I
       ia - IS-IS inter area, * - candidate default, U - per-user s
       o - ODR, P - periodic downloaded static route

Gateway of last resort is not set

    4.0.0.0/32 is subnetted, 1 subnets
O       4.4.4.4 [110/2] via 192.168.1.4, 00:00:11, FastEthernet0/1
C       192.168.1.0/24 is directly connected, FastEthernet0/1
```

```
Andheri#sh ip route
Codes: C - connected, S - static, R - RIP, M - mobile, B - BGP
       D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter
       N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external t
       E1 - OSPF external type 1, E2 - OSPF external type 2
       i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1, L2 - I
       ia - IS-IS inter area, * - candidate default, U - per-use
       o - ODR, P - periodic downloaded static route

Gateway of last resort is not set

    1.0.0.0/32 is subnetted, 1 subnets
C       1.1.1.1 is directly connected, Loopback0
    2.0.0.0/32 is subnetted, 1 subnets
O       2.2.2.2 [110/2] via 10.0.0.2, 00:28:18, FastEthernet0/0
    3.0.0.0/32 is subnetted, 1 subnets
O       3.3.3.3 [110/3] via 10.0.0.2, 00:27:14, FastEthernet0/0
    10.0.0.0/24 is subnetted, 2 subnets
C       10.0.0.0 is directly connected, FastEthernet0/0
O       10.0.1.0 [110/2] via 10.0.0.2, 00:27:24, FastEthernet0/0
```

```
Andheri#sh ip route vrf RED

Routing Table: RED
Codes: C - connected, S - static, R - RIP, M - mobile, B - BGP
       D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter are
       N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type
       E1 - OSPF external type 1, E2 - OSPF external type 2
       i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1, L2 - IS-IS
       ia - IS-IS inter area, * - candidate default, U - per-user st
       o - ODR, P - periodic downloaded static route

Gateway of last resort is not set

    4.0.0.0/32 is subnetted, 1 subnets
O       4.4.4.4 [110/2] via 192.168.1.4, 00:07:42, FastEthernet0/1
C       192.168.1.0/24 is directly connected, FastEthernet0/1
```

We now need to repeat this process for R3 & R6 Router 6 will peer OSPF using process number 2 to a VRF configured on R3. It will use the local site addressing to 192.168.2.0/24

```
Mahim(config)#INT L00
Mahim(config-if)#
*May 29 17:18:58.903: %LINEPROTO-5-UPDOWN: Line protocol is down
Mahim(config-if)#ip add 6.6.6.6 255.255.255.255
Mahim(config-if)#ip ospf 2 area 2
Mahim(config-if)#int f0/0
Mahim(config-if)#ip add 192.168.2.6 255.255.255.0
Mahim(config-if)#ip ospf 2 area 2
Mahim(config-if)#no shut
```

```
Churchgate(config)#int f0/1
Churchgate(config-if)#no shut
Churchgate(config-if)#ip add
*May 29 17:23:19.111: %LINK-3-UPDOWN: Interface FastEthernet0/1 is down
*May 29 17:23:20.111: %LINEPROTO-5-UPDOWN: Line protocol is down
Churchgate(config-if)#ip add 192.168.2.3 255.255.255.0
```

We also need to configure a VRF onto R3 as well.

```
Churchgate(config-if)#ip vrf RED
Churchgate(config-vrf)#rd 4:4
Churchgate(config-vrf)#route-target both 4:4
```

```
Churchgate(config-vrf)#int f0/1
Churchgate(config-if)#ip vrf forwarding RED
% Interface FastEthernet0/1 IP address 192.168.2.3 removed due to enabling VRF RED
Churchgate(config-if)#int f0/1
Churchgate(config-if)#ip add 192.168.2.1 255.255.255.0
```

```
Churchgate#sh run int f0/1
Building configuration...

Current configuration : 119 bytes
!
interface FastEthernet0/1
 ip vrf forwarding RED
 ip address 192.168.2.1 255.255.255.0
 duplex auto
 speed auto
end
```

Check the router in vrf RED

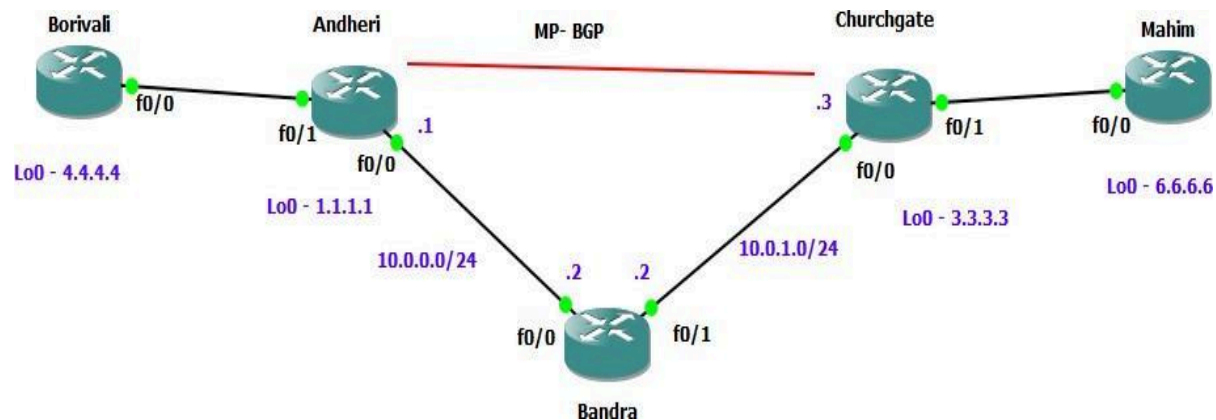
```
Churchgate#sh ip route vrf RED

Routing Table: RED
Codes: C - connected, S - static, R - RIP, M - mobile, B - BGP
       D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter area
       N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type 2
       E1 - OSPF external type 1, E2 - OSPF external type 2
       i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1, L2 - IS-IS level-2
       ia - IS-IS inter area, * - candidate default, U - per-user static route
       o - ODR, P - periodic downloaded static route

Gateway of last resort is not set

6.0.0.0/32 is subnetted, 1 subnets
O       6.6.6.6 [110/2] via 192.168.2.6, 00:01:10, FastEthernet0/1
C       192.168.2.0/24 is directly connected, FastEthernet0/1
```

Ok so we have come a long way now let's review the current situation. We now have this setup



```
Borivali#sh ip route
Codes: C - connected, S - static, R - RIP, M - mobile, B - BGP
       D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter area
       N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type 2
       E1 - OSPF external type 1, E2 - OSPF external type 2
       i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1, L2 - IS-IS level-2
       ia - IS-IS inter area, * - candidate default, U - per-user static route
       o - ODR, P - periodic downloaded static route

Gateway of last resort is not set

4.0.0.0/32 is subnetted, 1 subnets
C       4.4.4.4 is directly connected, Loopback0
C       192.168.1.0/24 is directly connected, FastEthernet0/0
```

As expected we have the local interface and the loopback address. When we are done we want to see 6.6.6.6 in there so we can ping across the MPLS. Check the routes on R1


```

Andheri#sh ip route
Codes: C - connected, S - static, R - RIP, M - mobile, B - BGP
       D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter
       N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external ty
       E1 - OSPF external type 1, E2 - OSPF external type 2
       i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1, L2 - IS
       ia - IS-IS inter area, * - candidate default, U - per-user
       o - ODR, P - periodic downloaded static route

Gateway of last resort is not set

    1.0.0.0/32 is subnetted, 1 subnets
C       1.1.1.1 is directly connected, Loopback0
    2.0.0.0/32 is subnetted, 1 subnets
O       2.2.2.2 [110/2] via 10.0.0.2, 00:28:18, FastEthernet0/0
    3.0.0.0/32 is subnetted, 1 subnets
O       3.3.3.3 [110/3] via 10.0.0.2, 00:27:14, FastEthernet0/0
    10.0.0.0/24 is subnetted, 2 subnets
C       10.0.0.0 is directly connected, FastEthernet0/0
O       10.0.1.0 [110/2] via 10.0.0.2, 00:27:24, FastEthernet0/0

```

```

Andheri#sh ip route vrf RED

Routing Table: RED
Codes: C - connected, S - static, R - RIP, M - mobile, B - BGP
       D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter area
       N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type 2
       E1 - OSPF external type 1, E2 - OSPF external type 2
       i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1, L2 - IS-IS
       ia - IS-IS inter area, * - candidate default, U - per-user sta
       o - ODR, P - periodic downloaded static route

Gateway of last resort is not set

    4.0.0.0/32 is subnetted, 1 subnets
O       4.4.4.4 [110/2] via 192.168.1.4, 00:07:42, FastEthernet0/1
C       192.168.1.0/24 is directly connected, FastEthernet0/1

```

```

Andheri(config)#router bgp 1
Andheri(config-router)#address-family ipv4 vrf RED
Andheri(config-router-af)#redistribute ospf 2
Andheri(config-router-af)#exit
Andheri(config-router)#end

```

```

Churchgate(config)#router bgp 1
Churchgate(config-router)#address-family ipv4 vrf RED
Churchgate(config-router-af)#redistribute ospf 2
Churchgate(config-router-af)#end

```



```

Andheri#sh ip bgp vpnv4 vrf RED
BGP table version is 9, local router ID is 1.1.1.1
Status codes: s suppressed, d damped, h history, * valid, > best, f
               r RIB-failure, S Stale
Origin codes: i - IGP, e - EGP, ? - incomplete

   Network          Next Hop          Metric LocPrf Weight Path
Route Distinguisher: 4:4 (default for vrf RED)
*> 4.4.4.4/32        192.168.1.4          2             32768 ?
*>i6.6.6.6/32        3.3.3.3              2          100      0 ?
*> 192.168.1.0       0.0.0.0              0             32768 ?
*>i192.168.2.0       3.3.3.3              0          100      0 ?

```

```

Churchgate#sh ip bgp vpnv4 vrf RED
BGP table version is 9, local router ID is 3.3.3.3
Status codes: s suppressed, d damped, h history, * valid, > best,
               r RIB-failure, S Stale
Origin codes: i - IGP, e - EGP, ? - incomplete

   Network          Next Hop          Metric LocPrf Weight Path
Route Distinguisher: 4:4 (default for vrf RED)
*>i4.4.4.4/32        1.1.1.1              2          100      0 ?
*> 6.6.6.6/32        192.168.2.6          2             32768 ?
*>i192.168.1.0       1.1.1.1              0          100      0 ?
*> 192.168.2.0       0.0.0.0              0             32768 ?

```

Which it is! 6.6.6.6 is now in the BGP table in VRF RED on R3 with a next hop of 192.168.2.6 (R6) and also 4.4.4 is in there as well with a next hop of 1.1.1.1 (which is the loopback of R1 – showing that it is going over the MPLS and R2 is not in the picture)

```

Andheri(config)#router ospf 2
Andheri(config-router)#redistribute bgp 1 subnets

```

```

Churchgate(config)#router ospf 2
Churchgate(config-router)#redistribute bgp 1 subnets

```

Before we do let's see what the routing table look like on R4

```

Borivali#sh ip route
Codes: C - connected, S - static, R - RIP, M - mobile, B - BGP
       D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter area
       N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type 2
       E1 - OSPF external type 1, E2 - OSPF external type 2
       i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1, L2 - IS-IS level-2
       ia - IS-IS inter area, * - candidate default, U - per-user static route
       o - ODR, P - periodic downloaded static route

Gateway of last resort is not set

  4.0.0.0/32 is subnetted, 1 subnets
C      4.4.4.4 is directly connected, Loopback0
  6.0.0.0/32 is subnetted, 1 subnets
O IA   6.6.6.6 [110/3] via 192.168.1.1, 00:00:50, FastEthernet0/0
C      192.168.1.0/24 is directly connected, FastEthernet0/0
O IA   192.168.2.0/24 [110/2] via 192.168.1.1, 00:00:50, FastEthernet0/0

```

Do the same step of on R6

```

Mahim#sh ip route
Codes: C - connected, S - static, R - RIP, M - mobile, B - BGP
       D - EIGRP, EX - EIGRP external, O - OSPF, IA - OSPF inter area
       N1 - OSPF NSSA external type 1, N2 - OSPF NSSA external type 2
       E1 - OSPF external type 1, E2 - OSPF external type 2
       i - IS-IS, su - IS-IS summary, L1 - IS-IS level-1, L2 - IS-IS level-2
       ia - IS-IS inter area, * - candidate default, U - per-user static route
       o - ODR, P - periodic downloaded static route

Gateway of last resort is not set

    4.0.0.0/32 is subnetted, 1 subnets
O IA   4.4.4.4 [110/3] via 192.168.2.1, 00:00:22, FastEthernet0/0
    6.0.0.0/32 is subnetted, 1 subnets
C       6.6.6.6 is directly connected, Loopback0
O IA   192.168.1.0/24 [110/2] via 192.168.2.1, 00:00:22, FastEthernet0/0
C     192.168.2.0/24 is directly connected, FastEthernet0/0

```

Lets chevk ping command

```

Borivali#ping 6.6.6.6
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 6.6.6.6, timeout is 2 seconds:
!!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 112/120/128 ms

```

Which we can – to prove this is going over the mpls and be label switched and not routed, lets do a trace

```

Borivali#trace 6.6.6.6
Type escape sequence to abort.
Tracing the route to 6.6.6.6

 0  192.168.1.1 20 msec 32 msec 24 msec
 1  10.0.0.2 [MPLS: Labels 17/19 Exp 0] 112 msec 136 msec 124 msec
 2  192.168.2.1 [MPLS: Label 19 Exp 0] 72 msec 92 msec 92 msec
 3  192.168.2.6 140 msec 124 msec 124 msec

```